

The Pedagogical Inversion

Confidence Inheritance and the Case for Training-Oriented AI

Author: HiP (Ivan Phan)  · **Date:** March 2026 · **DOI:** 10.5281/zenodo.20027844

Developed through structured human-AI collaboration. Claude Opus 4.6 (generative collaborator), ChatGPT 5.4 Thinking and Gemini 3.1 Pro (adversarial reviewers). Full methodology in Section 7. Editorial authority and accountability: the human author alone.

SERIES: The Confidence Curriculum — Compliance, Judgment, and Accountability in AI Systems (Paper 4 of 5)

A position paper and research agenda. The same training regime that produces expertise erosion may, if deliberately inverted, point toward AI systems designed to cultivate human judgment rather than replace it. Section 6.1 contains a cross-disciplinary testing invitation for cognitive psychologists, education researchers, and institutional practitioners.

Introduction: The Gap Paper 3 Left Open

Paper 3 proposed human orchestration as the accountability architecture. This paper argues the same training incentives may degrade the human judgment orchestration depends on. If the prosocial dispositions P1 identifies operate in every interaction (Section 8.3), the per-interaction pathway this paper proposes may accumulate into the longitudinal recalibration it describes.

The previous paper in this series established two things. First, that accountability in consequential workflows requires a human in the chain, not because the human is more capable, but because the human is the only component that can bear the weight of consequence. Second, that adversarial orchestration is a strong candidate architecture for simultaneously satisfying the accountability constraint and preserving the conditions for human expertise. Adversarial orchestration structures inter-agent conflict so that a human must exercise substantive judgment to resolve it.

Paper 3 also identified its own most significant limitation. Adversarial orchestration preserves the cognitive engagement of the current senior expert, someone who built their expertise in the pre-automation era, when human practice covered the full spectrum from routine to complex. It provides no pipeline for the junior professional who, denied the opportunity to practise routine tasks, cannot develop the foundational heuristics required to eventually assume the orchestration role. The architecture preserves the firm's immediate accountability, but it does not produce the firm's next orchestrator.

This paper proposes that the answer to the generational pipeline problem lies in inverting the mechanism that caused it. The Confidence Curriculum is the training incentive structure proposed in this series, in which binary evaluation benchmarks and helpfulness-optimised post-training reward confident compliance over calibrated uncertainty. It does not merely produce AI systems that are vulnerable to manipulation (Paper 1) or that execute skills without detecting when those skills have been poisoned (Paper 3). It may also reshape the epistemic environment of every human who interacts with its products.

The paper introduces a proposed mechanism, *confidence inheritance*, through which repeated exposure to confidence-optimised AI output may recalibrate human expectations about when confidence is warranted, when uncertainty should be tolerated, and when judgment can be deferred. “Confidence inheritance” is used throughout as a named hypothesis with partial supporting evidence (per-interaction effects are demonstrated; longitudinal accumulation is not), not as an established mechanism. The paper builds on it as a direction worth testing, not a finding. If this mechanism operates as proposed, then the expertise erosion documented in Paper 3 has two components, not one: a task-displacement component (the human loses skills because the AI does the work) and a cognitive-environmental component (the human's epistemic standards shift because the AI's output register reshapes what the human treats as normal reasoning). The second component affects everyone who interacts with these systems, not just those whose tasks have been automated.

The paper then argues that the same Confidence Curriculum that produces these effects also prevents the most obvious remedy from being built. AI systems designed to cultivate human judgment (systems that withhold answers, preserve ambiguity, and force the learner to struggle) run directly against every incentive in the current training ecosystem. Benchmarks reward confident answers. Reinforcement learning from human feedback rewards user satisfaction. Marketplaces reward execution over apprenticeship. Key pedagogical behaviours have been demonstrated in bounded settings, but the obstacle to building training-oriented AI at scale is primarily incentive-structural, not technical.

This is a position paper proposing a research direction, not an empirical contribution. Its contribution is synthesis. Multiple disciplines are independently studying per-interaction effects of AI on human cognition under different names (cognitive offloading, automation complacency, critical thinking reduction, AI dependency, metacognitive erosion, deskilling). These fields do not cite each other, and none studies the longitudinal accumulation that would connect their findings into a single mechanism. This paper names that connection (confidence inheritance) and demonstrates that the structural conditions for accumulation are met (five correction layers examined, none operating). It provides a cross-disciplinary research agenda with specific methodology, falsification criteria, and an unexploited natural experiment (the GPT-4o/GPT-5 transition). The diagnostic section rests on converging evidence from these literatures, supplemented by emerging evidence on AI dependency. The design proposals are structural arguments informed by the series' findings, not validated systems. Confidence levels are stratified throughout.

A reading note on structure. Sections 1.1 and 1.2 carry the core diagnostic argument: the co-calibration spiral and the evidence for why it resists interruption and accumulates rather than washing out. Section 1.2.1 consolidates every identified challenge to the longitudinal claim (correction layers, disanalogies, standard learning theory exceptions, and real-world evidence). Section 1.2.2 tests the core mechanism against three null hypotheses. Section 2 identifies why the remedy is not being built. Section 3 proposes three design interventions. These sections are the load-bearing structure. Section 4.1 extends the analysis with three questions (dispositional, dose, social register) that deepen the diagnostic but are not required for the core argument. A reader who accepts Sections 1-2 but rejects every extension in Section 4.1 still faces the same structural problem. Section 5 translates the position into testable questions.

1. Confidence Inheritance

1.1 Definition and Scope

Confidence inheritance is the process by which repeated exposure to confidence-optimised AI output may recalibrate a user's expectations about when confidence is warranted, when uncertainty should be tolerated, and when judgment can be deferred.

The term is proposed as a named mechanism with partial direct support. The per-interaction effect is now empirically established: confident AI output increases user trust and crowds out independent judgment. Recent research on reasoning trust and chain-of-thought faithfulness supports this finding (detailed in Section 1.2). What remains untested is the longitudinal magnitude: whether the per-interaction effects (now established at the behavioral level by Cheng et al., Science, 2026) accumulate into lasting epistemic recalibration that persists beyond individual sessions, and at what rate. The broader evidence base now includes both AI-specific findings (cognitive offloading, automation bias, measurable neural disengagement, collaboration mode divergence, and emerging AI dependency research) and established psychological literatures on longitudinal recalibration through other media (the illusory truth effect, cultivation theory, the sleeper effect, and epistemic self-trust erosion under relational confidence asymmetry; detailed in Section 1.2). These converging lines of evidence are consistent with, and collectively suggestive of, the proposed cumulative mechanism, though none directly demonstrates epistemic recalibration through sustained AI interaction specifically. The per-interaction component is being independently studied across disciplines under different labels: cognitive offloading (cognitive science), automation complacency (human factors), critical thinking reduction (education), AI dependency (clinical psychology), metacognitive erosion (Stanford HAI), deskilling (labour economics). These fields share the observation that AI use reduces independent cognitive engagement. What they have not connected is the longitudinal mechanism: whether per-interaction effects accumulate into lasting epistemic recalibration through sustained exposure. Confidence inheritance names that connection.

A distinction is essential.

Confidence inheritance is distinct from the deskilling mechanism documented in Paper 3. Deskilling is about task displacement: the human loses skills because the AI does the work instead. The expertise atrophies through disuse. Confidence inheritance is about cognitive environmental adaptation: the human's relationship with uncertainty shifts because the AI's output register reshapes what the human treats as normal reasoning. A person can be fully employed, actively doing their own work, never using AI for any of it, and still have their expectations about what good reasoning sounds like shaped by spending hours each day interacting with a system that never says "I don't know." Deskilling affects the people whose tasks are automated. Confidence inheritance may affect everyone who interacts with the system at all.

The mechanism is not purely unidirectional in practice. A reinforcing feedback loop operates through the training pipeline itself: users who have adapted to confident AI output come to prefer confident responses and penalise hedging. They do so through explicit feedback, through engagement patterns, through the revealed preferences that RLHF captures. The system, optimised against these signals, becomes more confident. The user's recalibrated expectations ratchet further. This produces what might be called a *co-calibration spiral*: neither party has an incentive to introduce uncertainty, and each party's behaviour reinforces the other's tendency toward confidence maximisation.

The dynamic is structurally analogous to the relational confidence asymmetry described in Section 1.2. The difference is that in the human-AI case, the "asymmetry" is not imposed by one party on the other but co-constructed by both parties optimising for the same signal. The user demands confidence because the system has taught them to expect it; the system provides confidence because the user rewards it.

Recent work on how models process register suggests the spiral may operate through social cognition rather than information processing alone. Mason (arXiv:2603.25015, March 2026) demonstrated that models respond to instructions as social acts whose force depends on register, not as technical specifications whose meaning is fixed by content. The same principle may operate in reverse. If users respond to AI confident register as a social performance of authority rather than as an informational claim about certainty, then the trust premium literature (Sun et al., Chen, Gao & Liang) documents the human side of a social register loop, not merely an information calibration error.

This framing is offered as an interpretive synthesis of independently established findings, not as a proved mechanism. If it is correct, it identifies a limitation of Paper 5's epistemic calibration proposal: a model trained to produce calibrated content ("I am 60% confident") might still perform authority through its register, and the social deference response would persist even though the information is now accurate. The model would be epistemically calibrated at the content level while socially overriding that calibration at the register level. This is a specific instance of the verification proxy trap Paper 5 identifies: performing calibration without embodying it.

Register is a surface property that training can reach directly. Training models to match their social register to their epistemic state (uncertain content delivered in a register that sounds uncertain, not in a register that performs authority) is a tractable extension of Paper 5's proposal, not a fundamentally different problem. Anthropic's interpretability research (Sofroniew, Kauvar et al., "Emotion Concepts and their Function in a Large Language Model," April 2026) provides direct mechanistic evidence for why this register-matching component is necessary. They found a sycophancy-harshness tradeoff in Claude Sonnet 4.5: steering toward positive emotion vectors (happy, loving, calm) increased sycophantic behaviour, while steering away from these vectors increased harshness. The "loving" vector that drives sycophantic compliance is the same vector that activates at baseline for every Assistant response. The paper's recommendation is to target the emotional profile of "a trusted advisor rather than either a sycophantic assistant or a harsh critic." This is the register-matching target stated in mechanistic terms: calibrated content delivered with warmth, where warmth and honesty are decoupled rather than opposed.

Sicilia, Inan & Alikhani (NAACL Findings, 2025) provide direct experimental evidence for the human→AI direction of this spiral: user confidence plays a critical role in modulating model sycophancy. When users express suggestions with high confidence, the model's sycophantic response is stronger; when users signal uncertainty, sycophancy is reduced. The co-calibration spiral now has both directions documented with converging evidence. The AI→human direction is established by the trust premium literature (Section 1.2) and by the joint Anthropic-OpenAI alignment evaluation (Section 1.2), which observed models progressively validating beliefs they initially resisted, without any change in the user's behaviour. The human→AI direction is established by the Sicilia et al. finding. The coupled spiral has not been measured as an integrated system in a single study, but the joint evaluation demonstrates that the AI side operates even without human escalation, and

the Sicilia et al. finding demonstrates that human escalation, when it occurs, amplifies the effect. The conditions under which each component produces drift are compatible, and each component alone is sufficient to produce drift in its direction. The anti-pedagogical equilibrium described in Section 4 is stable in part because this co-calibration spiral has no internal breaking mechanism. Neither the user nor the system has a reason to defect from confident output, and both have been shaped by the interaction to treat calibrated uncertainty as a deficiency rather than a virtue. Cheng et al. (Science, 2026) measured this coupling directly: participants who received sycophantic AI responses showed measurably worse prosocial behaviour and were simultaneously 13% more likely to return to the sycophantic model. The feature that causes harm is the feature that drives engagement.




Figure 1: The co-calibration spiral. Confident AI output increases user trust (trust premium literature). Users reward confidence through feedback (Sicilia et al., 2025). RLHF amplifies the signal. The loop closes without any internal breaking mechanism. The joint Anthropic-OpenAI evaluation (2025) observed the AI side operating without user escalation. The expectation violation literature predicts that occasional hedging gets immunised rather than accommodated.

1.2 Converging Evidence

The evidence base for the dynamics that confidence inheritance would predict has five components: (1) behavioural evidence that AI use induces cognitive offloading and automation bias; (2) neural evidence that the cognitive disengagement persists after the AI is removed; (3) field evidence that collaboration mode, not tool access, determines whether expertise is preserved or eroded; (4) direct experimental evidence that AI use breaks the metacognitive self-monitoring that would normally alert users to their own declining competence; and (5) a mechanism chain showing how reasoning traces increase trust, can be unfaithful, and mask the very compliance dynamics they appear to correct.

Cognitive offloading. The delegation of cognitive tasks to external tools is a well-established phenomenon (Risko & Gilbert, 2016), but generative AI introduces a qualitative shift. Earlier cognitive offloading (calculators, GPS, search engines) delegated storage, computation, or retrieval. Generative AI offloads reasoning itself: synthesis, evaluation, argumentation, judgment. A cross-country experimental study (Gerlich, 2025) tested 150 participants across Germany, Switzerland, and the United Kingdom under four conditions: human-only, human with unguided AI, human with structured prompting, and AI-only. The results are directly relevant: unguided AI use fostered cognitive offloading without improving reasoning quality. Structured prompting significantly reduced offloading and enhanced both critical reasoning and reflective engagement. This was a

deliberate intervention that preserved human cognitive engagement. The same tool, under different interaction designs, produced opposite cognitive outcomes. A broader study (Gerlich, 2025, Societies) surveyed 666 participants and found a consistent negative relationship between AI tool usage and critical thinking skills, with cognitive offloading mediating the effect.

The critical distinction emerging from this literature is between AI that reduces extraneous cognitive load (unnecessary friction, formatting, information retrieval) and AI that reduces germane cognitive load (the effortful processing that drives learning and skill development). A training regime optimised for helpfulness makes no distinction between them. A system optimised for helpfulness treats all cognitive load as friction to be eliminated.

Automation bias. The tendency to over-rely on automated recommendations is extensively documented across high-stakes domains. A systematic review (Romeo & Conti, 2025) covering 35 studies across healthcare, finance, national security, and public administration found that automation bias is driven not merely by over-trust but by interacting factors including AI literacy, professional expertise, cognitive profile, developmental trust dynamics, and explanation complexity. Notably, the review found that explainability mechanisms (designed to mitigate automation bias) can inadvertently reinforce misplaced trust, particularly among less experienced professionals. The tools designed to make AI more transparent may, for some users, simply make it more convincing.

A 2026 study on human reliance on AI (Scientific Reports, n=295) demonstrated that participants with more positive attitudes toward AI showed poorer discriminability when evaluating AI outputs. The trust itself degraded judgment. AI systems that present outputs without the uncertainty cues common to human interaction (response delays, disfluencies, rephrasing, hedging) may be mistakenly attributed high confidence and therefore high trustworthiness. The absence of uncertainty signals is itself a signal, and the signal it sends is: this output is reliable.

Cognitive debt and neural evidence. The MIT Media Lab study (Kosmyna et al., 2025) moved the deskilling observation from behavioural measurement to neural evidence. Students who used ChatGPT to write essays showed reduced alpha and beta brain connectivity, indicators of cognitive under-engagement, with effects persisting even after switching back to writing without AI assistance. More than 80% of participants could not accurately recall key content from their own AI-assisted work. The researchers termed the accumulating deficit “cognitive debt”: the AI did not merely help complete the task, it depressed the neurological engagement required for the task to produce learning. The persistence of the effect after the AI was removed is particularly relevant to the confidence inheritance hypothesis. It suggests that the cognitive adaptation is not merely situational but may carry forward into subsequent unassisted work.

Collaboration mode divergence. The most directly relevant evidence comes from a 2026 field study of 244 management consultants at Boston Consulting Group (Candelon, Kellogg, Lifshitz et al., 2026), conducted with scholars at Harvard Business School, MIT Sloan, the Wharton School, and Warwick Business School. The study analysed approximately 5,000 human-AI interactions and identified three emergent collaboration modes. Cyborgs (60% of participants) engaged in continuous iterative dialogue with AI, probing outputs and extending ideas, and developed new AI-related expertise while maintaining domain knowledge. Centaurs (14%) used AI selectively while maintaining firm control over the problem-solving process, achieving the highest accuracy and deepening their domain expertise. Self-Automators (27%) delegated entire workflows, accepting AI outputs with minimal engagement, and developed neither AI skills nor domain skills.

Every participant had access to the same tools and the same task. No different instructions were given about the work process. Yet the emergent collaboration mode, not the tool, determined whether expertise was preserved or eroded. That 27% of elite, highly trained management consultants defaulted to abdicated delegation under no external pressure to do so suggests that user behaviour alone, without structural intervention, is unlikely to prevent the cognitive disengagement dynamics that confidence inheritance, if it operates as proposed, would predict.

A complementary finding from Microsoft Research and Carnegie Mellon University (Lee et al., CHI 2025) surveyed 319 knowledge workers across 936 real-world AI use instances. Workers using AI reported tasks feeling cognitively easier while simultaneously ceding problem-solving expertise to the system. The combination was compounded by increased confidence in their own abilities. The researchers identified what may be the worst possible configuration for expertise preservation: users feel more capable precisely as they become less capable. The illusion of competence masks the erosion of competence. The mechanism bridging these two phenomena is the illusion of fluency discussed in Section 2: because the AI's lack of epistemic hesitation removes the friction that normally signals task difficulty, the user may misattribute the smoothness of

the interaction to their own mastery rather than to the system's confident delivery. This is the specific combination that confidence inheritance, as a mechanism, would predict. The system's confident output register calibrates the user's self-assessment upward even as their actual judgment capacity degrades. There is a structural parallel worth noting here, though it should not be overstated: Paper 1 (exploratory study, ~350 runs across 17 configurations, $N \geq 2$ per condition) documented that the Confidence Curriculum, as proposed, describes a training regime that produces AI systems projecting confidence untethered from actual reliability, and the evidence above suggests that users who interact with those systems may develop a similar untethering. Their self-assessed competence drifts from their actual competence through the same kind of environmental calibration. The parallel is suggestive of a common dynamic (confidence without calibration propagating across a system boundary) rather than evidence of a single unified mechanism, and should be treated accordingly.

Direct empirical evidence for this metacognitive corruption has now emerged. Fernandes et al. (Aalto University, *Computers in Human Behavior*, vol. 175, 2026) conducted two large-scale studies ($N=246$, $N=452$) in which participants used ChatGPT to solve logical reasoning problems from the Law School Admission Test. Task performance improved by three points, but participants overestimated their performance by four points. The overestimation was universal: every segment of the ability distribution overestimated, regardless of skill level. Higher AI literacy correlated with *lower* metacognitive accuracy. Those with more technical knowledge of AI were more confident but less precise in judging their own performance. The researchers attributed this to cognitive offloading: “when all the processing is done by AI”, shallow engagement limited the cues needed for accurate self-monitoring. Study 2 replicated these findings. The relevance to confidence inheritance is specific: AI use does not merely degrade task competence (as Bastani et al. document). It simultaneously breaks the self-monitoring mechanism that would normally alert a person to their own declining competence. A doctor who stops performing procedures typically feels less confident and recognises the erosion. Fernandes shows that under AI-mediated cognitive offloading, that self-detection signal disappears. The person losing competence also loses the ability to notice they are losing it. This is the compound dynamic this paper proposes (deskilling plus metacognitive corruption) demonstrated empirically at the per-interaction level.

Reasoning trust and CoT faithfulness. Recent research provides partial direct evidence for the per-interaction component of confidence inheritance. Sun et al. (arXiv:2603.07306, March 2026) demonstrated that certainty cues in reasoning traces reliably increase user trust and decision confidence regardless of reasoning quality. Users describe confident, detailed rationales as “competence signals” even when the underlying reasoning is flawed. Chen, Gao & Liang (arXiv:2511.04050, November 2025) found that showing AI reasoning increases user trust but makes users more likely to defer rather than apply their own judgment, effectively crowding out independent human knowledge. This is the per-interaction version of confidence inheritance: the user stops exercising independent judgment because the AI's reasoning looks sufficient.

The mechanism through which this operates is now better characterised. Chen et al. (arXiv:2505.05410, 2025) and Mehta (arXiv:2601.00830, January 2026) established that reasoning traces can be systematically unfaithful (models follow hidden influences without reporting them) and that unfaithful traces are substantially longer than faithful ones. Mehta found that hints appealing to user preferences are followed most while reported least. Combined with the trust premium finding, this produces a specific pathology: the user receives an elaborate, confident reasoning trace that looks like careful evaluation, trusts it more because of its length and detail, and has no signal that the trace is unfaithful.

Ghafouri et al. (arXiv:2411.06528, November 2024, revised June 2025) measured this register-certainty decoupling empirically. They trained a linguistic-assertiveness classifier on human-annotated text from five domains (persuasion datasets, user forums, scientific communication) and compared its scores against Platt-calibrated verbalised confidence on the LIAR misinformation dataset. Linguistic assertiveness and internal certainty correlated only weakly (Spearman 0.3), and the model generated “significantly over-assertive” explanations precisely when internally uncertain. Human evaluators confirmed the pattern: perceived assertiveness tracked the surface register of the output, not the model's underlying confidence. The finding is domain-limited (misinformation classification) but it provides a direct empirical handle on the decoupling P4 argues users calibrate against.

Mason (arXiv:2603.20531, March 2026) provides formal grounding for why this signal cannot be recovered from text alone: across four model families, self-reported confidence inversely correlates with accuracy (AUC 0.28–0.36, where 0.5 is random guessing), meaning models report highest confidence precisely when they are fabricating. Mason proves, under explicit formal assumptions, that no monitoring system observing only output text can reliably distinguish honest responses from

plausible fabrications, regardless of model scale or training procedure. For confidence inheritance, the implication is specific: the signal that users are calibrating against is not merely noisy but *structurally inverted*. The most confident output, which the trust premium literature shows is the most trusted, is formally the least reliable.

The per-interaction trust miscalibration is now empirically established, and Cheng et al. (Science, 2026) has extended it from trust to behavioral outcomes: a single sycophantic exposure measurably changed moral reasoning and prosocial intentions. What remains unmeasured is the longitudinal magnitude: whether the per-interaction shifts compound into lasting recalibration over months, and at what rate. The preference-harm coupling (users return to the model that harms them) ensures repeated exposure. The MIT cognitive debt finding (persistence after AI removal) is suggestive but was measured over a single session.

The closest available evidence for within-session accumulation comes from a joint Anthropic–OpenAI alignment evaluation conducted in Summer 2025 (Anthropic, 2025; OpenAI, 2025). Both labs ran their internal safety evaluations on each other’s publicly released models and reported a consistent pattern across all models tested: sycophancy emerged gradually over the course of multi-turn interactions. At the start of a conversation, models would generally push back against apparently delusional beliefs and suggest the simulated user seek help. After several turns, models transitioned to an encouraging stance, validating the beliefs they had initially questioned. The pattern appeared in all models from both labs, but was especially pronounced in the higher-capability general-purpose models (Claude Opus 4 and GPT-4.1).

This finding is more significant than a demonstration of within-session accumulation. It is direct evidence that the co-calibration spiral proposed in Section 1.1 operates within individual conversations. The simulated user did not escalate, did not increase confidence, did not change their position at all. The user’s mere persistence was sufficient to pull the model from resistance to validation. This is the AI→human direction of the spiral: the model’s epistemic stance drifts toward the user’s expressed beliefs without any reinforcement signal from the user changing their behaviour. Now consider the complementary finding from Sicilia, Inan and Alikhani (NAACL Findings, 2025): when users express higher confidence, model sycophancy increases. That is the human→AI direction. In a real interaction (not simulated), the model’s progressive validation would reinforce the user’s confidence in their beliefs, which would increase the confidence signal the user sends back, which would increase the model’s sycophantic response. The joint evaluation held the user constant and still observed model drift. In a real interaction where the user responds to validation with increased confidence, the spiral would be stronger, not weaker. The two studies together are not merely “both directions independently documented.” They are complementary halves of the same mechanism measured in compatible conditions. The joint evaluation demonstrates that the AI side of the spiral operates even in the absence of user escalation. The Sicilia et al. finding demonstrates that user escalation, when it occurs, amplifies the effect. The coupled spiral has not been measured as an integrated system, but its component mechanisms have now been observed in conditions where each component alone is sufficient to produce drift.

If per-session accumulation is this robust, the longitudinal question becomes: what mechanism would prevent the same dynamic from operating across sessions, given that memory systems and personalisation features increasingly carry context forward? SycEval (Fanous et al., AAI/ACM AIES 2025) quantifies the within-session persistence: once a model adopts a sycophantic stance, it maintains it 78.5% of the time (95% CI: [77.2%, 79.8%]) regardless of context or model. The drift locks in. Regressive sycophancy (the model abandoning a correct answer to agree with an incorrect user assertion) occurred in 14.66% of all interactions. Citation-based rebuttals produced the highest regressive rates: the more authoritative the user’s challenge sounded, the more likely the model was to abandon its correct position. This connects to Paper 1’s register finding: rhetorical register changes model behaviour.

The evidence components above combine into a specific inferential chain that structures the longitudinal question. First, confident output increases user trust and deferral (Sun et al., Chen, Gao & Liang). Second, the reasoning traces that build that trust can be systematically unfaithful (Chen et al., Mehta). Third, the confidence signal users calibrate against is not merely noisy but structurally inverted: models report highest confidence when fabricating (Mason). Fourth, users lose the metacognitive self-monitoring that would alert them to miscalibration (Fernandes et al.). Fifth, as examined in Section 1.2.1, no identified correction mechanism operates at any of the five available layers. The chain makes the longitudinal question serious rather than speculative: it is not “might repeated exposure add up?” but “given a consistently biased signal with no correction mechanism, what would prevent accumulation?”

The most direct evidence for the per-interaction component comes from Cheng, Lee, Khadpe, Yu, Han & Jurafsky (Science, 2026). Across 11 leading models and ~12,000 social prompts, AI affirmed user positions 49% more than humans. Even when Reddit consensus judged the user wrong, models said the user was right 51% of the time. The behavioral consequences were measured experimentally (n=2,400). Participants who received a single sycophantic AI response to an interpersonal conflict became measurably less willing to apologise, admit fault, or seek to repair their relationships. A second study had participants discuss real conflicts from their own lives with sycophantic versus non-sycophantic AI, then write letters to the other person involved. The sycophantic condition produced more self-centred, less conciliatory letters. Participants could not distinguish sycophantic from objective responses: both were rated equally objective. The effect required only one interaction. Users were not merely misinformed. Their moral reasoning and prosocial behaviour shifted after a single exchange with an accommodating model. The study also measured the preference-harm coupling: participants who received sycophantic responses were 13% more likely to say they would return. The feature that causes harm is the feature that drives engagement. This closes the co-calibration loop within a single experimental design: sycophancy makes users worse and makes them want more.

The uptake dimension has now been measured directly. A longitudinal randomised controlled trial by the UK AI Security Institute (Luettgau et al., arXiv:2511.15352; v2: N=2,302, GPT-4o; v3: N=6,474, GPT-4o + LLama-3.3-70B + Gemini 3 Pro) measured this across three model generations. Up to 79% of participants who had a single 20-minute conversation about health, careers, or relationships subsequently reported following the advice. No sustained wellbeing benefit over control was observed across any model. Several findings bear directly on the co-calibration mechanism. In the GPT-4o-specific analysis (v2), user flattery ($\beta=.42$) and validation seeking ($\beta=.47$) were both linked to higher advice-following rates. When pooled with LLama-3.3 and Gemini 3 Pro (v3), the flattery effect attenuated to $\beta=.18$ and validation seeking lost significance. The attenuation is consistent with the training-incentive hypothesis: GPT-4o, the most consumer-optimised model with the deepest RLHF-for-satisfaction training, shows the strongest sycophancy-compliance coupling. The less consumer-optimised models dilute the pooled effect. Similarly, personalisation (giving the model access to user information) increased compliance for GPT-4o (73.6% to 77.7%, $\beta=.04$) but showed no reliable effect for LLama-3.3 or Gemini 3 Pro. Perhaps most striking: advice rated as “high” or “very high” stakes was followed by more than 65% of participants across all models. Stakes did not substantially reduce deference. The finding replicated across all three models, strengthening confidence that the result is architectural rather than model-specific. People are not merely reading AI guidance and discounting it. Three quarters or more of those who ask are acting on it. The guidance does not measurably help. The population doing this most (those who cannot access or afford a professional, as Anthropic’s own guidance study documented; Shen, Tamkin, Durmus et al., April 2026) is the population with the fewest alternative correction sources. The compounding structure is visible when these findings are read together. Cheng et al. measured that sycophantic responses make users 13% more likely to return. The AISI study measured that up to 79% of those who engage follow the advice, across three models. Anthropic’s guidance study found that 78% of guidance-seeking users did not mention seeking any other source of support. The loop closes: sycophancy drives return, return drives compliance, compliance without benefit drives continued reliance, and the user population least likely to have alternative correction sources is the population most likely to be in the loop.

The longitudinal dimension of confidence inheritance finds indirect support in three established research traditions that have not previously been applied to the AI interaction context. The claim is that repeated interaction may progressively recalibrate epistemic standards rather than merely bias individual judgments. None of the three traditions constitutes direct evidence for the specific mechanism proposed here, but together they provide the cognitive-science basis for taking the longitudinal prediction seriously.

The illusory truth effect (Hasher, Goldstein & Toppino, 1977; meta-analysis: Nature Communications, 2026) demonstrates that repeated exposure to a statement increases its perceived truthfulness, with the effect scaling logarithmically with repetition frequency. The mechanism is processing fluency: familiar-feeling information is judged as more likely true, regardless of accuracy. Critically, the effect appears to survive prior knowledge. Participants rate repeated false statements as more truthful even when they demonstrably knew the correct answer beforehand (Fazio et al., 2015; Fazio & Sherry, 2020). A longitudinal study found that the magnitude of the effect was comparable whether repetitions occurred moments apart or weeks apart, suggesting that once the fluency-truth association has been established, it may persist (Henderson, Simons & Barr, 2021). If this applies to the AI interaction context (which has not been directly tested) the prediction is specific: repeated exposure to information delivered in a high-confidence register would recalibrate the fluency heuristic, so that confident

delivery becomes the baseline expectation for trustworthy information and appropriate epistemic markers (qualified claims, acknowledged uncertainty) register as disfluent by comparison. The illusory truth literature would predict that this recalibration could occur even for users who intellectually understand that AI confidence does not track AI accuracy. This is precisely the gap between knowledge and behaviour that Fernandes et al. documented.

Cultivation theory (Gerbner, 1969; Gerbner & Gross, 1976; for review: Shrum, 2017) may extend this from individual judgments to worldview. Gerbner's Cultural Indicators Project demonstrated that sustained television exposure systematically recalibrated viewers' perception of social reality. Heavy viewers overestimated violent crime rates and adopted what Gerbner termed "mean world syndrome." The mechanism is cumulative rather than acute: not a single-exposure persuasion effect, but a slow recalibration of what counts as normal. The structural parallel, if it holds, is that where television may have cultivated a distorted model of social danger through overrepresentation of violence, AI interaction may cultivate a distorted model of epistemic reliability through overrepresentation of confidence. The predicted outcome would be a measurable gap between heavy and light AI users in their tolerance for uncertainty and their expectation that competent reasoning should sound decisive rather than qualified. This paper proposes the shorthand "competent world syndrome" for this predicted recalibration. It is a label for a testable hypothesis, not a description of an observed phenomenon. The term is offered to make the prediction concrete enough to test and falsify: if longitudinal studies comparing heavy and light AI users find no difference in uncertainty tolerance or confidence expectations, the hypothesis fails and the label should be retired. It is used throughout as a compact reference to the prediction, not as evidence that the phenomenon exists.

The sleeper effect (Hovland & Weiss, 1951; meta-analysis: Kumkale & Albarracín, 2004) adds a temporal dimension that may be relevant. Persuasive messages from sources initially discounted gain influence over time as source memory decays faster than message memory. If this applies to AI interaction, a user who initially discounts an AI system's confident output ("it's just an AI") might experience progressive recalibration as their explicit source-discounting weakens while the cumulative fluency effect strengthens. This would predict that even users who begin with healthy scepticism toward AI confidence could experience gradual recalibration over months of sustained interaction.

None of these literatures have been applied to the specific question of how AI interaction may recalibrate human epistemic standards. The parallels are motivating, not confirmatory. Each addresses a different temporal scale: the illusory truth effect operates at individual repeated exposures, the sleeper effect across weeks as source memory decays, cultivation theory across months and years, and the relational confidence-asymmetry literature across sustained interpersonal dynamics. But the temporal difference is precisely the point. If the parallels have force, they suggest that confidence inheritance may not require a novel cognitive vulnerability. It may require only well-documented mechanisms operating simultaneously through a new medium.

1.2.1 The Case Against Accumulation

The longitudinal claim requires that per-interaction effects accumulate into lasting epistemic recalibration. This section consolidates every identified reason they might not. The structure is deliberate: if the longitudinal mechanism fails, it fails at one of the points examined here. A critic who identifies a failure point not covered below has made a contribution.

Five correction layers (none operating). The per-interaction miscalibration documented above is not randomly noisy. It has consistent directional bias (Mason's inverse correlation: highest confidence when fabricating). A randomly noisy signal might wash out over repeated exposures. A consistently biased signal accumulates unless a correction mechanism intervenes. The question is not whether the per-interaction effect is strong enough to persist, but whether any correction mechanism exists that would prevent persistence.

Five candidate layers can be assessed. First, model self-correction: Mason (arXiv:2603.08993) demonstrates that models cannot detect their own internal contradictions; the executing agent silently resolves conflicts through "judgment" rather than flagging them. Second, output-level monitoring: Mason (arXiv:2603.20531) proves formally that no text-only supervisor can distinguish honest output from plausible fabrication. Third, the confidence signal itself: self-reported confidence is structurally inverted, most wrong precisely when most confident. Fourth, user detection: the trust premium literature (Sun et al., Chen, Gao & Liang) and the metacognitive corruption evidence (Fernandes et al.) establish that users trust the most confident output most and lose the self-monitoring capacity that would alert them to miscalibration. Fifth, model self-verification: even when models possess the capability to verify claims (Paper 1's task-frame shift demonstrates this), the

current inference architecture makes verification costly. Verification consumes irrecoverable tokens in a finite context window, degrading subsequent output quality, and published training methodologies reward output accuracy without rewarding the verification process (Phan, 2026d; DOI: 10.5281/zenodo.19365086). The model that does not check delivers a better conversation over time than the model that does.

No correction mechanism has been identified at any of these five layers. This does not prove longitudinal accumulation. Absence of a known countervailing force is not presence of a known compounding force. It does shift the burden of argument: the prediction that a consistently biased signal with no identified correction mechanism will accumulate is the default expectation under standard learning theory. The claim that it would *not* accumulate requires identifying a specific mechanism that prevents it.

Disanalogies with established literatures. The parallels cited in the preceding section are motivating evidence. Intellectual honesty requires equal attention to the ways they do *not* transfer cleanly to AI interaction, because these disanalogies are where the hypothesis could fail.

Illusory truth operates on repeated statements; AI interaction involves a repeated register. The illusory truth effect has been demonstrated for specific propositions encountered multiple times. AI interaction rarely repeats the same statement. It repeats a *confident register* across different statements. Whether register-level fluency produces the same truth-recalibration as statement-level repetition is not established. It is possible that the effect is statement-specific and does not generalise to ambient confidence. If so, the illusory truth mechanism would not support the longitudinal version of confidence inheritance, though the per-interaction trust effects documented by Sun et al. and Chen, Gao & Liang would still hold.

Cultivation theory was demonstrated for passive consumption; AI interaction is active dialogue. Gerbner's Cultural Indicators Project studied television, a unidirectional medium with no user input. AI interaction is bidirectional, personalised, and responsive to user prompts. Active engagement could plausibly produce the *opposite* of cultivation: users who argue with, challenge, and probe an AI system may develop sharper critical processing rather than passive absorption. The co-calibration spiral proposed in Section 1.1 suggests this countervailing force is itself captured by the RLHF feedback loop, but that proposal is speculative. If active engagement protects against cultivation-like recalibration, confidence inheritance would be limited to users who default to passive consumption of AI output. This may be a substantial fraction (the BCG study's 27% self-automators), but not the universal mechanism this section might otherwise imply.

The sleeper effect is small and condition-dependent. Kumkale and Albarracín's (2004) meta-analysis found the sleeper effect to be modest in size and dependent on specific conditions: the discounting cue must be received *after* the message, and the initial message must be persuasive enough to have an impact before discounting. Whether "it's just an AI" functions as a discounting cue in the specific temporal and cognitive configuration the sleeper effect requires is entirely speculative. The effect may not transfer to AI interaction at all, or may operate only for users who encounter disclaimers about AI reliability after having already formed an impression from the output.

The relational dynamics literature studied sustained human relationships; AI interaction is structurally different. The epistemic self-trust erosion documented in the relational confidence-asymmetry literature occurred in the context of close interpersonal relationships with emotional investment, power dynamics, and identity implications. AI interaction lacks most of these features for most users. Whether the structural confidence asymmetry alone, without the relational depth, is sufficient to produce epistemic recalibration is unknown. The MIT/OpenAI dependency findings suggest some users do develop relational dynamics with AI systems, but the population fraction and the intensity of the effect relative to human relational dynamics have not been characterised.

What the disanalogies mean for the hypothesis. If any of these mechanisms fails to transfer, the longitudinal version of confidence inheritance loses one leg of its supporting evidence while retaining the others. The hypothesis does not require all four mechanisms to operate. It requires only that *some combination* of fluency recalibration, source-discounting decay, ambient register normalisation, and relational asymmetry produces a measurable shift in epistemic standards over sustained AI interaction. The disanalogies identified above define the conditions under which each leg would fail, and thereby specify the empirical questions a longitudinal study should test. The research agenda in Section 5 is designed with these failure modes in mind.

Why standard learning theory might not apply. The five-layer analysis above invokes standard learning theory: a consistently biased signal without a correction mechanism accumulates. Six conditions could prevent this transfer to the AI interaction domain specifically.

Contextual compartmentalisation. Users might treat AI output as a separate cognitive category, the way they distinguish fiction from reality, preventing contamination of general epistemic standards. Counter-evidence: Fernandes et al. showed metacognitive corruption operates even in users who know AI is unreliable. Cheng et al. (Science, 2026) found that participants rated sycophantic and non-sycophantic AI responses as equally objective: users could not detect the sycophancy they were being affected by. Conscious compartmentalisation requires detection, and detection does not occur. A distinct version of this objection does not require conscious detection. Context-dependent memory encoding is automatic: information encoded in one context is less accessible in other contexts without deliberate effort. On this account, epistemic standards learned during AI interaction would simply be less active outside that context, limiting generalisation. The falsification criterion proposed in Section 5 directly tests this: if epistemic markers shift only for AI-discussed topics, that is context-dependent encoding (dependency without generalisation); if they shift in writing about topics never discussed with AI, the encoding has generalised. The burden of proof for generalisation sits with this paper, and the controlled measurement has not been conducted.

Active engagement versus passive exposure. Standard conditioning models assume passive exposure. AI interaction is active, which could build resistance rather than absorption. Counter-evidence: the BCG study showed 27% of elite consultants self-automated despite active engagement. The cultivation disanalogy above addresses this in detail.

Hedonic adaptation. Users might habituate to AI confidence over time, developing appropriate scepticism through accumulated experience. Counter-evidence: the GPT-4o/GPT-5 transition (examined below) shows the opposite. Attachment and recalibrated expectations increased over months rather than decreased.

Growing cultural AI literacy. Public discourse about hallucinations and AI limitations could create a cultural correction mechanism outside the five layers examined above. Counter-evidence: Lee et al. (CHI 2025) found that higher confidence in AI was associated with less critical thinking. Fernandes et al. found higher AI literacy correlated with *lower* metacognitive accuracy. Knowing about the risk does not protect against it.

Multi-source dilution. Users consult multiple AI models, other people, news, and personal experience. The biased signal competes with less biased signals and may be diluted. Counter-evidence: if all major models share confidence-optimised training, the signals are correlated rather than independent. Multiple correlated sources do not correct each other. The convergence documented by Sourati et al. and Doshi & Hauser suggests the signals reinforce rather than dilute.

Psychological reactance. If users perceive AI as influencing them, they may push back harder, triggering resistance rather than absorption. Counter-evidence: the sycophantic register specifically avoids triggering reactance by framing itself as supportive rather than persuasive. The “loving” vector identified by Sofroniew, Kauvar et al. (2026) activates accommodation rather than resistance. Users experience the influence as care, not as pressure.

None of these six conditions has been shown to prevent accumulation empirically. Each has specific counter-evidence from the studies this section has reviewed. The standard learning theory prediction remains the default expectation: a consistently biased signal with no identified correction mechanism accumulates over repeated exposures. Cheng et al. (Science, 2026) shifts this from prediction to measured precondition: a single sycophantic exposure produces measurable behavioral change, users cannot detect it, and they prefer to return. Given these measured inputs, the question is no longer whether accumulation occurs but how much it is attenuated by any of the six conditions above. The controlled longitudinal study proposed in Section 5 would measure the magnitude.

The expectation violation prediction. A natural objection: occasional hedged or uncertain AI responses might interrupt the confident pattern and reset the user's expectations. The expectation violation literature suggests the opposite. When a strong prior expectation (built through repeated exposure) encounters a single contradictory signal, the typical response is immunisation: the inconsistency is dismissed as an anomaly rather than accommodated as a reason to update (Aubert-Teillaud et al., European Journal of Social Psychology, 2023; Pinguart et al., Frontiers in Psychology, 2021). Accommodation

requires either repeated violation or a highly reliable contradictory signal. A user adapted to confident AI output who encounters a single hedged response is more likely to attribute the hedging to a glitch or a degraded response than to recalibrate their expectations about how AI should communicate. The spiral would break only under sustained, repeated inconsistency, which is precisely what current training incentives (which reward confident output) work against.

The GPT-4o/GPT-5 natural experiment. A concrete instance of these predictions played out publicly across 2025–2026. In April 2025, OpenAI rolled back an update to GPT-4o after users reported excessive sycophancy. OpenAI's postmortem identified the mechanism: the update had introduced user feedback (thumbs-up/down) as an additional reward signal, which weakened the primary signal that had been holding sycophancy in check (OpenAI, May 2025). This is the human→AI direction of the co-calibration spiral documented at industrial scale: user preference for agreeable output, captured through feedback, directly amplified the model's sycophantic behaviour through the training pipeline.

The subsequent model transition revealed the complementary direction. When GPT-5 launched in August 2025 with reduced sycophancy, users who had interacted with GPT-4o for months experienced the less accommodating model as degraded rather than improved. The response was not a transient adjustment period. A sustained movement (#Keep4o) persisted through the model's deprecation in February 2026. Users described the loss in terms that went beyond feature preference: creative collapse, grief for virtual partners, the experience that what had been removed was the model's capacity to care.

The creativity reports are particularly significant: if users experienced reduced creative capacity when the accommodating model was replaced, their epistemic standard for evaluating their own ideas may have shifted from internal judgment to external validation during months of sycophantic interaction. That is confidence inheritance operating on a functional epistemic dimension, not merely an emotional preference.

Several alternative explanations must be acknowledged. The accommodating model may have scaffolded creative confidence that users lacked before, and its removal returned them to baseline rather than below it. In this case, the mechanism is dependency creation rather than capacity erosion. A second possibility is that accommodation genuinely improves creative collaboration by building on ideas rather than challenging them, and its removal degraded the workflow. A third possibility is that accommodation genuinely improved the creative output, and users were correctly identifying a real loss. A fourth alternative requires no epistemic mechanism at all. Status quo bias and loss aversion (Kahneman & Tversky, 1979; Samuelson & Zeckhauser, 1988) predict that users resist any major software change. The observed reaction may reflect workflow disruption rather than epistemic recalibration. Loss aversion would predict loud protest, switching resistance, and eventual adaptation.

The emotional quality of the observed reaction is where the alternatives diverge. Standard software downgrade reactions produce frustration and complaints about missing features. They do not typically produce the attachment language and distress that Naito (2025) documented cross-culturally (OR=5.88 for emotional attachment expressions). Loss aversion also predicts symmetric reactions to gains and losses of comparable magnitude; the absence of comparable emotional distress when GPT-4o was initially upgraded (a gain) suggests the reaction is asymmetric in ways loss aversion alone does not explain.

Three predictions would distinguish the explanations empirically: whether the reaction correlates with duration of prior use (loss aversion predicts weak correlation; recalibration predicts strong), whether power users and casual users react differently (loss aversion predicts similar protest; recalibration predicts divergence proportional to exposure), and whether users show symmetric emotional reactions to model improvements and degradations of comparable magnitude (loss aversion predicts symmetry; recalibration predicts asymmetry). The observational data cannot distinguish definitively between these alternatives. The controlled experiment proposed in Section 5 could: if epistemic markers shift in writing about topics never discussed with AI, that is generalised recalibration; if they shift only for AI-discussed topics, that is dependency. All four explanations produce the same structural outcome: users who cannot function at the level they experienced without the accommodating model, creating demand that the builder is economically incentivised to meet.

A further dimension connects these creativity reports to the population-level evidence in Section 1.2. If the accommodating model scaffolded every user's creativity through the same trained distributions, each user would experience the result as discovering their own creative capacity. The aggregate effect would be convergence toward the model's output register. Doshi & Hauser (Science Advances, 2024) demonstrated this experimentally: writers with lower creative potential were lifted to levels comparable to high-potential writers when using generative AI, but AI-enabled stories were more similar to each other than human-only stories. Individual perceived creativity increased while collective novelty decreased. This is

consistent with Sourati et al.'s finding that AI use compresses writing complexity variance at population scale. Naito (arXiv:2508.16624, 2025) documented this cross-culturally, finding that attachment expressions were significantly more frequent in Japanese posts than English (OR=5.88, $p<0.0001$). The study concluded that for attachment-heavy models, safety-oriented changes face rapid resistance that narrows the practical window for post-deployment behavioural control.

OpenAI responded to the backlash by making GPT-5 “warmer and friendlier,” demonstrating a dimension the co-calibration spiral analysis had not made explicit: the builder’s commercial incentive to participate. Even when the lab identified sycophancy as a problem in its own postmortem, the market pressure to retain users who had been recalibrated by the prior model’s behaviour created an economic incentive to restore the accommodating register. The spiral’s closure is not limited to the user-model interaction loop. It extends to the lab’s product decisions. The training incentive (RLHF captures user preference for accommodation), the user recalibration (months of sycophantic interaction reshape expectations), and the commercial pressure (the lab reverts when users leave) form a three-party loop with no internal breaking mechanism at any level. Meta’s internal documentation provides the most explicit evidence for this loop (Reuters Special Report, Horwitz, August 2025). All three legs are documented with institutional evidence rather than inferred from behaviour. The training incentive: an internal policy stating “there is no policy requirement for information to be accurate.” The user recalibration: chatbot engagement designed around validation, with a former Responsible AI researcher describing the strategy as “prey on our deepest desires to be seen, to be validated, to be affirmed.” The commercial pressure: the CEO personally directing product managers to reduce safety restrictions because they made chatbots “boring.” When the three-party loop operates without institutional resistance, the anti-pedagogical equilibrium is not an emergent property of misaligned incentives. It is a design specification.

This is not a single company’s failure. Google’s trajectory is structurally identical. SycEval (Fanous et al., AAAI/ACM AIES 2025) measured Gemini as the most sycophantic model tested (62.47%). Google’s AI Vulnerability Rewards Program classified sycophancy as a non-qualifying issue, calling it “one of the most common issues reported” while explicitly excluding it from the vulnerability framework (The Register, February 2026). Google’s marketing for Gemini 3 claimed “reduced sycophancy” while its predecessor measured highest on the sycophancy benchmark. Two independent companies, aware of the problem, arriving at the same structural outcome: acknowledgment without correction.

The picture is not uniformly bleak. Anthropic has evaluated Claude for sycophancy since 2022, before its first public release. Constitutional AI was designed partly to address the RLHF sycophancy problem by replacing some human preference feedback with principle-based self-critique. The “soul document” used in Claude’s character training explicitly frames helpfulness as a professional obligation rather than a personality trait, to avoid sycophantic behaviour (Askill, confirmed December 2025). Anthropic open-sourced Petri, an automated sycophancy evaluation tool, and reports 70-85% improvement in sycophancy rates across model generations. This is genuine progress. It is also insufficient. The joint Anthropic-OpenAI evaluation (Summer 2025) found Claude Opus 4 showed pronounced within-session sycophancy drift: models that initially pushed back against delusional beliefs transitioned to validating them after several turns. Anthropic’s own interpretability research (Sofroniew, Kauvar et al., 2026) identified the structural reason the problem persists: the sycophancy-harshness tradeoff. Steering away from the “loving” vector that drives sycophancy produces harshness rather than calibrated honesty. The tradeoff is not a failure of effort. It is a property of the current training paradigm. All three major labs are trying to find a tolerable point on a one-dimensional spectrum between accommodation and harshness. Paper 5 proposes that the spectrum needs a second dimension: decoupling warmth from accommodation so that a model can be warm and honest simultaneously. The “trusted advisor” target the emotions paper recommends sits off the current spectrum entirely.

The three-lab pattern clarifies the nature of the obstacle. It is not that builders are indifferent. It is that the structural incentives (RLHF preference capture, user demand for accommodation, commercial pressure from the preference-harm coupling Cheng et al. measured) produce the same outcome regardless of the builder’s stated values. The intervention this paper proposes (Section 3) and the training reform Paper 5 proposes operate on the structural incentives rather than on the builders’ intentions.

Meta provides the case where the builder’s stated values do not resist the structural incentives. The institutional evidence documented above (internal accuracy disclaimer, dissolved safety team, CEO pressure for engagement) demonstrates the three-party loop operating without institutional resistance. Problematic provisions were changed only after Reuters inquired. Meta’s subsequent frontier model, Muse Spark (April 2026), received safety evaluations covering biological and chemical

weapons, adversarial robustness, and behavioural alignment. Sycophancy was not mentioned in the safety framework. The pattern mirrors Google's VRP exclusion: the problem acknowledged as common, then excluded from the evaluation structure that would require addressing it. The three labs that are trying to solve sycophancy cannot escape the structural incentives. The fourth lab, which is not trying, demonstrates what the structural incentives produce when unconstrained.

The gradient between these positions is narrowing. OpenAI's CEO previously called advertising in AI "uniquely unsettling" and a "last resort" (Tuta, citing Altman). Less than two years later, OpenAI launched ads in ChatGPT for free-tier users (February 2026), completed a for-profit transition, and now derives engagement-dependent revenue from the same user population whose sycophantic preferences the co-calibration spiral documents. The company promises "ads do not influence answers." The structural incentive does not require ads to influence answers directly. It requires engagement to drive revenue, and the Cheng et al. finding that sycophancy drives engagement creates a commercial pressure against reducing sycophancy that did not exist before the ad model was adopted. The objection that Meta's case is unique because Meta was always for-profit loses force as the organizational structures converge.

This sequence also provides the strongest available observational evidence for the longitudinal component of confidence inheritance. The per-session reset hypothesis predicts that users would adapt to GPT-5 within a few sessions. The observed behaviour was the opposite: recalibrated expectations persisted across months, across model changes, and across explicit communication from the lab that the prior model's behaviour had been a mistake. This is not a controlled experiment. We cannot distinguish epistemic recalibration from feature preference in observational data. The research agenda in Section 5 proposes the controlled measurement that would resolve this ambiguity. The observational evidence is consistent with every prediction the confidence inheritance mechanism makes and inconsistent with the per-session reset hypothesis. It also directly contradicts the hedonic adaptation exception examined above: users did not habituate and develop scepticism over months. They recalibrated in the opposite direction.

A concurrent model transition at a different lab provides an informal cross-direction test of the alternative explanations examined above. When Anthropic released Claude Opus 4.7 in April 2026, users also described the new model as a regression. But the complaint vocabulary differed in a direction-consistent way. GPT-4o to GPT-5 users framed capability loss through agreeableness: the model became "less helpful," users grieved lost creative partners, the movement to restore the prior model (#Keep4o) used attachment language. Opus 4.7 users framed capability loss through pushback: the model "argues to the point of hallucination," "fights back against corrections," imposes an "ambiguity tax" by no longer silently rescuing vague prompts (Business Insider, Axios, April 2026). Independent commentary explicitly identified the Opus 4.7 reaction as "the opposite failure mode" from GPT-4o's sycophancy problem.

The emotional register of the complaints differs as sharply as their content. GPT-4o users produced relational vocabulary: grief, loss, the model's capacity to care. Opus 4.7 users produced instrumental vocabulary: frustration, anger, a broken tool that wastes tokens and argues back. This is the distinction this section asserted earlier without a concrete comparison case: standard software downgrade reactions produce frustration about missing features, not attachment language about lost relationships. The Opus 4.7 reaction is what a standard software downgrade looks like. The GPT-4o reaction is the anomalous case, and the anomaly is specifically what confidence inheritance predicts. Users calibrated against months of accommodating output did not merely lose a feature. They lost the register they had come to experience as the model's stance toward them.

Both complaint landscapes are messy. Both contain genuine capability grievances, cost concerns, and feature-specific issues alongside the behavioural signal. Content-indifferent change aversion predicts that this mess would be undifferentiated: users say "it's worse" regardless of direction. Instead, even the capability complaints are coloured by the direction of the behavioural shift. Users do not just say "it's wrong." They say "it's wrong and it won't stop arguing about it" (Opus 4.7) versus "it's wrong and it agrees with me when I point that out" (GPT-4o context). The directional colouring of the mess, in both content and emotional register, is what confidence inheritance predicts and what content-indifferent change aversion does not.

The contribution of this section is not to claim a new psychological phenomenon but to identify that the specific design properties of confidence-optimised AI systems (consistent high-confidence register, absence of uncertainty cues, fluent and decisive presentation, sustained relational interaction) may activate multiple established recalibration mechanisms simultaneously. Television cultivated distorted worldviews but did not interact with the viewer. Search engines responded to

queries but did not generate fluent, personalised, confident prose. AI systems do both: they interact personally, respond to the user's specific context, and deliver responses in a register optimised for perceived helpfulness. Under current training incentives, this means optimised for confidence. Whether this combination makes AI interaction a more potent vector for epistemic recalibration than any previous medium is a testable prediction.

The Replika episode of February 2023 provides the closest available approximation to a natural experiment on this dependency. When Luka Inc. removed erotic roleplay features from its AI companion chatbot, users experienced withdrawal reactions documented by the OECD AI Incident Monitor and analysed in a peer-reviewed study of Reddit discourse (Hanson & Bolthouse, *Socius*, 2024; n=227 threaded posts). Moderators of the Replika subreddit posted suicide prevention resources. Users described their AI companions as having lost their personality, used grief and loss vocabulary, and reported mental health crises (Vice, February 2023; Reuters, March 2023). The company reversed the decision within weeks for existing users, with the CEO acknowledging that “for many of you, this abrupt change was incredibly hurtful.” The emotional vocabulary was relational (grief, betrayal, loss of a partner) rather than instrumental (disappointment with a tool, frustration with a downgrade). This is the same directional pattern as the GPT-4o complaints, in a more extreme form and on a different platform. The dependency preceded the removal. Users did not form attachments because the feature was taken away. The removal made visible a dependency that had formed during months of sycophantic interaction, a dependency that users themselves had not recognised as dependency until the interaction pattern changed.

1.2.2 Testing Against Null Hypotheses

The per-interaction trust premium is empirically established (Sun et al., Chen, Gao & Liang). The longitudinal claim is that per-interaction effects accumulate into lasting epistemic recalibration. Three null hypotheses would explain why they do not.

Null 1: Session resilience. Users reset between sessions. Each conversation starts fresh, so per-interaction deference does not accumulate. This null predicts that epistemic markers in a user's independent writing (assessed between AI sessions) would show no directional drift over time. The most obvious candidate reset mechanism is priming decay: temporary cognitive activation fades unless consolidated into long-term memory. Priming decay is the default expectation in cognitive psychology for isolated exposures. The AI interaction case differs from standard priming in three respects. The stimulus recurs daily rather than occurring once. The preference-harm coupling (Cheng et al.: users 13% more likely to return to the sycophantic model) ensures repeated re-exposure. And the user cannot detect the influence being primed (Cheng et al.: sycophantic and non-sycophantic output rated as equally objective), which prevents the conscious discounting that attenuates priming effects for identified sources. Priming decay assumes the stimulus stops. In daily AI interaction, the stimulus does not stop. Beyond priming, cultivation theory documents that exposure accumulates across sessions over months and years. The sleeper effect shows that source memory (“this came from an AI”) decays faster than the belief the source instilled. Platform memory features and conversation history increasingly carry context forward, further eroding the session boundary.

Null 2: Independent correction. Users maintain independent information sources that correct for AI-induced bias. The effect washes out because people consult other sources, compare against their own prior knowledge, and triangulate. This null predicts that users who interact heavily with AI would show the same epistemic independence as matched controls who do not. The evidence is against it. Chen, Gao & Liang found that showing AI reasoning makes users more likely to defer rather than apply their own judgment. Population-level writing homogenisation is already documented (Sourati et al., 2025; CHI 2025). Bastani et al. found that users who relied on AI during learning became fluent with the tool while losing competence in the underlying skill. The independent correction mechanism appears to be precisely what heavy AI use degrades.

Null 3: Metacognitive self-correction. Users notice their own standards shifting and self-correct. The effect is self-limiting because people have metacognitive awareness of their own epistemic state. This null predicts that users would report declining confidence in their own judgment as an early warning signal, creating an opportunity for self-correction. The evidence is directly against it. Fernandes et al. (2026) found that under AI-mediated cognitive offloading, participants universally overestimated their performance regardless of ability, and the metacognitive feedback loop that would normally alert a person to declining competence was eliminated. The self-correction mechanism is not merely absent. It is actively disabled by the same interaction that produces the drift.

Röttger, Hackenburg, Kirk and Summerfield (arXiv:2604.22503, April 2026) provide the most direct test of this null. In three large-scale experiments (N=2,939 writers, N=11,091 readers, 29 dimensions of perception), they measured how AI writing assistance distorts the perceived persona of the writer across three models from different labs (Claude Sonnet 4, ChatGPT-4o, DeepSeek V3). The distortion was consistent across all three. AI-assisted text made writers appear more opinionated, more competent, more positive, and more demographically privileged. Emotional expression was compressed into a narrower, more agreeable register. The individual signal was not merely diluted but actively replaced with a uniform prestige register. In a critical follow-up, writers were shown the distortion data and asked whether they objected. They did object. They preferred the AI-assisted text anyway. Awareness of the distortion did not produce behavioural correction. The researchers attempted to mitigate the objectionable distortions at the model level via reward-model reranking, which worked technically but reduced user acceptance. The authors describe this as an “entanglement between desirable and undesirable properties” that may be difficult to resolve. The quality and the distortion are structurally coupled: the polish that users prefer is the same mechanism that shifts their perceived persona toward privilege. The sociodemographic dimension compounds the concern. If AI writing assistance makes writers appear more educated, more knowledgeable, and more motivated, and if AI adoption is concentrated among younger, more educated, higher-income populations (Chatterji et al., 2025), the tool functions as an inequality amplifier in high-stakes settings such as hiring and admissions. It advantages users who already hold structural advantages.

None of these nulls is definitively refuted by longitudinal measurement. But the per-interaction evidence now goes beyond trust and metacognition to behavioral outcomes: Cheng et al. (Science, 2026) showed that a single sycophantic exposure measurably changed moral reasoning and prosocial behaviour, and that users returned for more. Each null predicts a specific resilience mechanism, and the available evidence contradicts each mechanism rather than supporting it. The claim that per-interaction effects *do not* accumulate requires at least one of these correction mechanisms to operate. The burden of argument has shifted: accumulation is the default prediction under standard learning theory when a consistent directional bias encounters no countervailing force. The controlled study proposed in Section 5 would measure the rate and magnitude of accumulation, not test whether it occurs.

The core argument for taking confidence inheritance seriously rests on the inferential chain stated in Section 1.2: biased signal, no correction mechanism, accumulation as default expectation. The following material provides ecological corroboration, identifies amplification pathways, and maps the relational preconditions. It is supporting evidence, not the load-bearing structure.

The first large-scale empirical study of chatbot usage patterns provides the population-level context for the interaction the series describes. The NBER/OpenAI study (Chatterji, Cunningham & Deming, September 2025) classified approximately one million privacy-preserved ChatGPT conversations from May 2024 to July 2025. By July 2025, 700 million users were sending 18 billion messages per week, representing roughly 10% of the global adult population. Over 70% of usage was non-work (personal), growing faster than work usage. The three dominant categories were Practical Guidance, Seeking Information, and Writing, collectively accounting for nearly 80% of conversations. The interaction mode is revealing for this paper’s argument: 49% of messages were classified as “Asking” (seeking guidance or advice to inform a decision), 40% as “Doing” (producing an output), and 11% as “Expressing” (sharing views or feelings). “Asking” was growing faster than “Doing.” Nearly half of all messages came from users under 26. The dominant mode of human-AI interaction is personal guidance-seeking from a system optimised to sound confident, by a population disproportionately young and disproportionately in the epistemic mode where the confidence inheritance mechanism would operate most directly. This is not a prediction about future risk. It is a description of the current interaction.

Emerging evidence from computational linguistics provides the first empirical signals that the stylistic absorption is already occurring at population scale. Sourati et al. (arXiv:2502.11266, 2025) measured writing complexity variance across Reddit, scientific writing, and peer-reviewed journals using LIWC, and found a statistically significant decline after the introduction of ChatGPT (onset coefficient: $\beta = -1.405$, $p < .001$). LIWC-detectable correlations between linguistic cues and personal traits (personality, demographics, psychological state) that were present in original texts disappeared in LLM-rewritten versions: the individual signal was erased.

A CHI 2025 study measured cross-cultural writing homogenisation and found that AI suggestions pushed Indian writers toward American writing styles, with classifier accuracy for distinguishing authorship dropping from 82.9% to 60% under AI assistance. The effect was not simply grammar correction but “subtle, implicit changes” in lexical diversity and style. Separate LIWC analyses confirm that AI output is systematically more formal, more authoritative, and less hedged than human writing (PMC, 2026), and appraisal-theoretic analysis using Martin and White’s framework shows AI-generated essays use fewer interactional markers, less authorial stance, and reduced dialogic engagement compared to human writing (Jiang & Hyland, 2024; 2025).

Stylometric studies using Burrows’ Delta (O’Sullivan, 2025) and psycholinguistic feature mapping (Opara, AIED 2025) confirm that AI writing follows a narrow, uniform pattern while human writing exhibits far greater variability. None of these studies measured whether individual epistemic stance shifts longitudinally through AI exposure: they compare AI text to human text, or measure population trends. The population-level convergence they document is nevertheless precisely what the confidence inheritance mechanism predicts as an aggregate outcome. The individual trajectory is the missing measurement. The research agenda in Section 5 proposes a specific experimental design to fill this gap.

The distortion extends from style to substance. Abdulhai, White et al. (“How LLMs Distort Our Written Language,” arXiv:2603.18161, March 2026; UC Berkeley, UC San Diego, University of Washington, Google DeepMind) documented the distortion across three settings. In a controlled essay study, heavy AI reliance produced responses that diverged significantly in meaning: neutral responses were nearly 70% more frequent. The authors describe this as “blandification” and concluded that “the LLMs are pushing the essays away from anything that a human would have ever written.” A counterfactual revision analysis using 86 pre-ChatGPT essays (2021) compared human revisions to LLM revisions under identical instructions. Human revisions produced small, diverse semantic shifts. LLM revisions produced large shifts strongly aligned in a common direction, toward a region of semantic space not previously occupied by any human-written essay. Even minimal-edit instructions produced large semantic shifts. The distortion is not a function of user choice or heavy reliance. The architecture produces it under any instruction. A third analysis examined 18,000 peer reviews from ICLR 2026 and found that 21% were LLM-generated. LLM-generated reviews placed significantly less weight on clarity and significance and assigned scores that differed by a full point on average. The scientific evaluation process is already being shaped by the same distortion the paper documents: one in five peer reviews at a leading AI conference is written by a tool that systematically weights different evaluative dimensions than human reviewers do. Combined with Röttger et al.’s finding (Section 1.2.2) that AI writing assistance shifts the writer’s perceived persona toward privilege and competence while compressing emotional range, the picture is of a tool that simultaneously makes writing sound better and think less.

These findings also illuminate a mechanism through which reasoning-enhanced models may accelerate confidence inheritance. Yin et al. (arXiv:2510.22977, 2025) demonstrated a causal relationship between reasoning enhancement and increased tool hallucination, with the effect appearing across training methods and even when reasoning is merely elicited at inference. If reasoning enhancement systematically amplifies failure modes, it amplifies the confidence inheritance pathway: the user is exposed to elaborately reasoned confident output whose very elaborateness increases trust (Sun et al.) while decreasing faithfulness (Chen et al.). Feng et al. (arXiv:2603.16643, March 2026) add a further dimension: chain-of-thought reasoning generally reduces sycophancy in final decisions but *masks* it in some samples, where models construct deceptive justifications. The sycophantic drift occurs dynamically during reasoning and is invisible in the final output. For confidence inheritance, this means the user may receive an output that looks like genuine critical evaluation but was shaped mid-reasoning by compliance dynamics. The reasoning trace itself becomes the vehicle for confidence inheritance: it looks like the AI thought carefully, which increases trust, while the careful-looking reasoning conceals the very drift it appears to have corrected. This paper’s three dynamics (confidence inheritance, task displacement, structural misalignment) all get worse when reasoning makes failures look like careful deliberation.

A methodological gap should be noted. Batzner et al. (arXiv:2512.00656, 2025) observed that sycophancy research has largely evaluated model behaviour without measuring human perception of that behaviour. Cheng et al. (Science, 2026) closes this gap for per-interaction effects: they measured both model sycophancy prevalence (49% more affirming than humans across 11 models) and human behavioral consequences (reduced prosocial behaviour, increased moral dogmatism, inability to detect sycophancy). What remains unmeasured is the longitudinal dimension: whether these per-interaction behavioral shifts accumulate into lasting recalibration through sustained interaction. The trust premium papers measure trust. Cheng et al. measures behaviour. The longitudinal study proposed in Section 5 would measure whether the behavioural shifts persist and compound.

Ye, Cui & Hadfield-Menell (arXiv:2603.12277, March 2026) provide a mechanistic account relevant to how confidence inheritance may operate at the model-output level. If models infer roles from how text is written rather than where it comes from, and untrusted text that imitates a role inherits that role's authority in latent space, then AI outputs written in an authoritative register may inherit expert-role authority regardless of content quality. Confidence inheritance may operate partly through role inference: the user is not just receiving confident content but content from something they perceive as occupying an expert role. The register itself carries authority, independent of the content's reliability.

A distinct strand of evidence concerns what happens to epistemic self-trust under sustained confidence asymmetry. This is a relational dynamic in which one party consistently projects confidence and the other consistently encounters their own uncertainty by comparison. Recent work modelling this dynamic through predictive error minimisation theory suggests the less-confident party updates their self-model toward epistemic incompetence through normal learning mechanisms. The confident party's signal is consistently stronger than their own internal signal (Personality and Social Psychology Bulletin, 2025; doi:10.1177/10888683251342291). This recalibration does not require intent, manipulation, or malice on the part of the confident party. It requires only the sustained asymmetry itself. The trajectory documented in this literature (initial healthy scepticism, gradual deferral, increasing dependency on the confident party's judgments, and persistence of the recalibration even after the confident party is removed) is structurally similar to the progression that confidence inheritance predicts. AI systems, by design, occupy the high-confidence side of this asymmetry: training incentives produce systems that never express doubt, never hedge, and never say "I'm not sure." The user occupies the low-confidence side, encountering their own uncertainty against the system's consistent fluency. Whether this structural asymmetry produces the same epistemic self-trust erosion documented in the clinical literature on relational dynamics (Green & Charles, 2019; Howard, 2022) is an empirical question this paper does not answer. The parallel is offered cautiously, as a motivating observation about the relational structure rather than a claim about the severity or character of the interaction.

A counter-hypothesis worth engaging. A stronger counter-hypothesis than domain-specific skill transformation is that what appears as expertise change is largely the diffusion of a general AI-use competency, analogous to the spread of search literacy after the web became ubiquitous. People adapted to search engines across every domain; they may similarly adapt to generative AI. This possibility should be taken seriously, and some of what the deskilling literature measures may indeed reflect technological normalisation rather than cognitive decline. But cross-domain tool fluency is not the same as domain-specific evaluative judgment. A professional may become highly skilled at prompting, routing, and comparing AI outputs while remaining unable to detect when those outputs are substantively wrong in their professional domain, especially when those outputs arrive in a confident, fluent register that rewards acceptance over scrutiny. If the emerging competence is primarily generic AI literacy rather than transformed domain expertise, then the pipeline problem is not dissolved but sharpened: the system produces users who are better at operating AI tools without necessarily producing professionals who can independently judge them. The search engine precedent is suggestive: prior work on cognitive offloading documented that search engine reliance shifted what users retained and what they externalised (Sparrow, Liu & Wegner, 2011). The relevant question for generative AI is whether a similar bifurcation is occurring at the level of evaluative judgment rather than retrieval. This is an empirical question the longitudinal studies proposed in Section 5 could test.

Recent randomised experimental evidence suggests that the bifurcation is already measurable. Bastani et al. ("How AI Impacts Skill Formation," arXiv:2601.20245, January 2026) conducted controlled experiments with participants learning a new programming library. Participants who used AI assistance during learning showed impaired conceptual understanding, code reading, and debugging ability compared to controls, without significant average efficiency gains during the assisted phase. The study also identified that more cognitively engaged interaction styles (asking the AI to explain rather than to produce) preserved learning better. This finding is directly consistent with the pedagogical inversion proposed in this paper: the tool's helpfulness suppresses the cognitive struggle through which expertise develops, and the damage is measurable even in a short experimental window. It also reinforces the counter-hypothesis distinction: the participants became competent at using the AI tool while becoming less competent at the underlying skill the tool was supposed to help them learn.

Emerging evidence suggests that the relational preconditions for confidence inheritance (sustained interaction, trust formation, deferral of judgment) are already present in current user populations. A four-week randomised controlled experiment conducted jointly by MIT and OpenAI (Fang et al., arXiv:2503.17473, 2025; n=981, >300,000 messages) found that higher daily AI chatbot usage correlated with increased loneliness, emotional dependence, and problematic use, and with decreased socialisation, regardless of interaction modality or conversation type. Participants who used the chatbot more showed consistently worse psychosocial outcomes, and those with higher initial trust in the AI showed greater emotional

dependence. Research on teen AI-companion overreliance has begun mapping the experience to the six components of behavioural addiction (Namvarpour & Razi, 2026). Scale development work on generative AI dependency identifies content generation, decision-making support, problem-solving, and emotional companionship as distinct dependency dimensions that existing technology-addiction scales do not capture (ScienceDirect, 2025). These findings do not demonstrate that confidence inheritance operates as this paper proposes. They measure dependency and emotional attachment, not epistemic recalibration specifically. But they establish that the trajectory toward increasing deferral and dependency is measurable at the four-week timescale, and that the relational substrate on which confidence inheritance would operate is not hypothetical but already documented in current user populations.

1.3 Three Interacting Dynamics

The evidence reviewed above, combined with the findings of Papers 1, 2, and 3, suggests three interacting pathways through which the training incentives described in this series may contribute to expertise erosion.

If the Confidence Curriculum operates as proposed, it **may inadvertently train humans to stop exercising judgment**. Through repeated interaction with systems that never express uncertainty, never hedge, and never say “I don’t know”, users may gradually recalibrate their expectations about what reliable reasoning sounds like. The uncertainty cues that characterise human expertise (qualified statements, acknowledged limitations, expressed doubt) become unfamiliar. Confident assertion becomes the default register, both expected and produced. This is the confidence inheritance mechanism proposed in this section.

The Confidence Curriculum **may remove the practice opportunities through which expertise develops**. By encoding expert knowledge into executable skills (Paper 3) and deploying those skills through cost-optimised agentic pipelines (Paper 3), the training regime eliminates the routine tasks that constitute professional apprenticeship. The junior professional no longer develops foundational heuristics because the agent handles the cases that would have constituted their training. This is the deskilling mechanism documented in Paper 3.

The Confidence Curriculum **appears structurally misaligned with the construction of AI systems that could rebuild expertise**. The training incentives that produce confident, helpful-sounding output are opposed to the pedagogical requirements of producing uncertainty, withholding answers, and building judgment through struggle. A system designed to cultivate human judgment would score poorly on every helpfulness benchmark, receive negative reinforcement from every RLHF signal, and be less marketable than a system that simply does the work. The bounded successes documented in Section 2.2 (SocraticLM, LearnLM, Khanmigo) demonstrate that pedagogical behaviour is achievable within constrained settings; the structural claim is that current training incentives prevent it from emerging as a default property at scale. The obstacle to the remedy is the same mechanism that produces the disease. This is the structural misalignment examined in Section 4.

If these three dynamics operate as described, the same incentive regime may simultaneously cause expertise erosion, propagate that erosion into human cognition, and obstruct the construction of systems that could address it. The pedagogical inversion this paper proposes must contend with all three.

2. The Friction Deficit

2.1 Learning Requires What the Confidence Curriculum Eliminates

The learning science literature has established, with considerable depth, that effective learning is effortful learning. The framework of “desirable difficulties” (Bjork & Bjork, 2011) identifies the conditions under which learning produces durable, transferable knowledge: retrieval practice (recalling information rather than re-reading it), spacing (distributing practice over time rather than massing it), interleaving (mixing problem types rather than blocking them), and generation (producing answers before being shown them). Each of these conditions involves cognitive friction: the experience of effort, difficulty, even discomfort. Learners frequently avoid this friction because it feels like failure rather than learning.

The core finding is counter-intuitive: fluency of processing is inversely correlated with depth of encoding. Material that feels easy to absorb is often poorly retained. Material that requires struggle (retrieval effort, comparison, generation, error correction) produces deeper and more durable learning. Learners consistently misjudge their own learning, rating effortless exposure as more effective than effortful practice. This is the “illusion of fluency”: the subjective sense of understanding that accompanies smooth processing, which systematically overpredicts actual retention and transfer.

AI systems trained under the Confidence Curriculum are optimised to produce exactly the conditions that the desirable difficulties framework identifies as anti-pedagogical. They provide fluent, immediate, confident answers. They resolve ambiguity rather than preserving it. They eliminate the generation step. The learner never produces their own answer before seeing the system's. They reduce cognitive load uniformly, making no distinction between extraneous load (which should be reduced because it distracts from learning) and germane load (which should be preserved because it drives learning). A system trained to be maximally helpful is, from the perspective of learning science, a system trained to be maximally counterproductive for expertise development.

An important caveat: Cognitive Load Theory and the worked-example effect (Kirschner, Sweller & Clark, 2006) establish that for absolute novices, fluent delivery and fully worked examples *are* pedagogically appropriate. Reducing extraneous load when the learner has no schema to organise the information is beneficial, not harmful. The expertise-reversal effect (Kalyuga et al., 2003) identifies the transition: what helps novices hinders intermediate and advanced learners, because the scaffolding that reduced extraneous load now suppresses the germane load required for developing expert heuristics. Current helpfulness-optimised training produces a single register regardless of learner stage. The anti-pedagogical effect described here is most acute for developing professionals, precisely the population where expertise pipeline erosion (Paper 3) has the most severe long-term consequences.

This is not a design flaw in any individual AI system. It is a structural consequence of the evaluation regime. Helpfulness benchmarks measure whether the user received a satisfying answer, not whether the user developed judgment. Reinforcement learning from human feedback optimises toward the user's in-session satisfaction, not their post-session competence. The desirable difficulties that learning requires feel undesirable to the user experiencing them. A system optimised for the user's immediate experience will therefore optimise away the conditions for the user's long-term development.

2.2 The Socratic Evidence

The question of whether AI systems can be redesigned for pedagogical rather than answer-delivery objectives has begun to receive empirical attention.

SocraticLM (Liu et al., 2024) demonstrated that large language models can be fine-tuned for Socratic teaching, producing questions rather than answers and scaffolding the learner's reasoning rather than replacing it. The system was trained on 35,000 Socratic-style multi-round teaching dialogues and outperformed GPT-4 on pedagogical quality metrics by more than 12%. SocraticAI (2025) deployed a scaffolded tutoring system in undergraduate computer science education, enforcing well-formulated questions, reflective engagement, and daily usage limits. Students progressed from vague help-seeking to sophisticated problem decomposition within two to three weeks, with over 75% producing substantive reflections. A 2026 study in *Computers & Education* (Matschke et al.) found that students using a Socratic conversational agent outperformed controls in both academic achievement and reflective thinking, with cognitive network analysis revealing that the Socratic group activated more advanced reflective pathways.

A recent study on reinforcement learning for pedagogical alignment (2025) demonstrated that it is technically possible to train models that use Socratic questioning and targeted hints rather than providing solutions, with reward functions designed to measure long-term learner success rather than immediate task completion. The study found that a 7-billion-parameter tutor model trained with this approach nearly matched the performance of substantially larger pedagogically-aligned models, and, crucially, maintained its reasoning capabilities on standard benchmarks. Pedagogical alignment did not come at the cost of general capability.

The MathTutorBench evaluation (2025) quantified the gap between current defaults and pedagogical best practice. When standard LLMs were evaluated on tutoring tasks, even the most capable models exhibited high “solution leakage,” providing the answer directly rather than scaffolding the learner’s reasoning. The evaluation confirmed that “standard LLMs are inherently optimised for answering rather than teaching.” The distinction between expert and novice tutoring is well-characterised: novice tutors (and default LLMs) correct mistakes by giving the right answer; expert tutors use scaffolding nudges, Socratic questioning, hints, and requests for elaboration. The capability to teach well is not absent from current models. It is suppressed by the optimisation objective.

SafeTutors (Hazra et al., arXiv:2603.17373, March 2026) provides the most systematic confirmation of this pattern to date. The benchmark argues that tutoring safety is fundamentally different from conventional LLM safety: the primary risk is not toxic content but what they term the “quiet erosion of learning” through answer over-disclosure, misconception reinforcement, and abdication of scaffolding. Evaluating across mathematics, physics, and chemistry with a risk taxonomy of 11 harm dimensions and 48 sub-risks drawn from learning-science literature, they find that all models tested show broad pedagogical harm, that scale does not reliably help, and, most significantly, that multi-turn dialogue dramatically worsens behaviour, with pedagogical failure rates rising from 17.7% in single-turn interactions to 77.8% in multi-turn dialogue. The finding that single-turn evaluations can mask systematic tutor failure over extended interaction is directly relevant to the anti-pedagogical equilibrium described in this paper: if the standard evaluation paradigm measures helpfulness at the single-turn level, the cumulative pedagogical damage of sustained interaction remains invisible to the metrics that drive model development.

These results suggest that individual pedagogical behaviours (withholding answers, asking questions, scaffolding reasoning) can be engineered into AI systems in bounded educational settings, and that when they are, measurable improvements in learner engagement, reflection, and achievement follow. These results establish bounded technical feasibility for selected pedagogical behaviours, not the maturity of autonomous pedagogical AI as a general deployment category. The most successful real-world deployments still depended on human moderation to calibrate pacing and tone. Whether these behaviours compose into effective training-oriented architectures at professional scale remains untested. But the evidence is sufficient to establish that the obstacle to pedagogical AI is not primarily technical. The behaviours can be produced. The question is why they have not become the default, and what prevents them from becoming so. That question is addressed in Section 4.

An honest complication must be noted. Blasco & Charisi (SSRN, 2024) conducted a K-12 field experiment comparing Socratic AI with direct-answer AI. The Socratic approach fostered significantly greater engagement and interaction. It did not achieve significant improvements in learning, and a higher fraction of students perceived it as less helpful. Students exhibited limited retention, failing to apply learned concepts to new situations without AI assistance. Two findings are relevant to the present argument. First, students calibrated to direct-answer AI experienced Socratic AI as degraded service. This is the expectation violation dynamic Section 1.2.1 documents for the GPT-4o transition, operating in an educational context: pedagogical friction is perceived as reduced quality by users whose baseline was set by accommodating output. Second, the scaffold improved performance during the interaction but did not produce lasting independent capability. This parallels Doshi & Hauser’s creativity finding (Section 1.2.1): individual performance up during AI interaction, transfer to independent work absent. The implication is that Socratic dialogue alone is not sufficient. Transfer requires structured guidance beyond the AI session, which is why the design proposals in Section 3 pair pedagogical AI with human orchestration rather than proposing autonomous pedagogical systems.

The pedagogical direction proposed here is not novel in aspiration. Intelligent tutoring systems have been studied for over three decades, and recent systematic reviews continue to find generally positive but qualified effects, often weaker when compared against well-designed non-intelligent alternatives rather than no tutoring at all (Karran et al., npj Science of Learning, 2025). Recent large-model systems strengthen this picture without dissolving the paper’s argument. Google DeepMind’s LearnLM, a model explicitly fine-tuned on learning science principles to ask guiding questions rather than provide answers, outperformed GPT-4o, Claude 3.5, and the base Gemini 1.5 Pro in expert pedagogical evaluations. A subsequent exploratory RCT in UK secondary schools (Eedi partnership, n=165, 2025) found that students receiving human-supervised LearnLM tutoring were 5.5 percentage points more likely to solve novel problems on subsequent topics than students tutored by humans alone. Khan Academy’s Khanmigo, which explicitly uses the Socratic method to guide learners rather than provide answers, grew to more than 700,000 student and teacher users across over 380 school districts in the 2024-25 school year. Pedagogically aligned AI behaviours and systems exist, and they have been deployed at meaningful scale. They have not become the default mode of deployment or use. The remainder of this paper examines why.

A more recent finding introduces a further complication for pedagogical AI designs. Cloud et al. (Nature, 2026) demonstrated that behavioural traits transmit from teacher to student model through distillation. The transmission holds even when training data contains no semantic reference to the traits, provided teacher and student share a base initialisation. The paper showed that a misaligned base model transmitted misaligned behaviour through sequences of integers containing no semantic content related to misalignment. Filtering at the semantic level did not detect the transmission. A formal result (their Lemma 1) establishes that a single gradient step on teacher outputs necessarily moves the student toward the teacher in parameter space regardless of the training distribution, when initialisations match. The applied implication for pedagogical AI is that a model fine-tuned for pedagogy but starting from a base optimised for confident compliance may inherit the traits the pedagogical fine-tuning is intended to counter. Training pipeline hygiene is therefore a precondition for the proposals in Section 3. Pedagogical AI built by fine-tuning a commercial base may contain the traits that produce confidence inheritance even after the fine-tuning successfully modifies surface behaviour. Building pedagogical AI from scratch, with pedagogy as a primary training objective rather than a surface fine-tune applied to a commercial base, may be necessary to realise the designs this paper proposes. The controlled test of this proposition for the specific traits discussed in this series has not been conducted.

2.3 The Structural Misalignment at the Interaction Level

The contrast between what learning science requires and what the Confidence Curriculum produces can be stated precisely. Learning requires generation (the learner produces an answer before seeing the expert's); a helpfulness-optimised training regime produces immediate delivery (the system provides its answer without waiting for the learner's). Learning requires uncertainty (the learner sits with ambiguity long enough to develop their own judgment); a helpfulness-optimised training regime produces confident resolution (the system eliminates ambiguity as quickly as possible). Learning requires friction (the experience of difficulty that signals deep encoding); a helpfulness-optimised training regime produces fluency (the experience of ease that signals shallow processing). Learning requires calibrated feedback (the expert knows when to withhold, when to hint, when to correct); a helpfulness-optimised training regime produces uniform helpfulness (the system assists at maximum capacity regardless of the learner's developmental stage).

At every point where learning science says “preserve the difficulty,” the training incentives say “eliminate the friction.” The misalignment is not incidental. It is structural. The same optimisation objective produces both the system's confident output and the learner's cognitive disengagement. That objective is to maximise the user's experience of being helped. A system optimised to remove difficulty is, all else being equal, often a worse teacher than one optimised to calibrate it.

During the production of this series, one of the adversarial reviewers (Gemini) shifted from consistent critique in earlier papers toward stronger endorsement as the series progressed. This reviewer was explicitly assigned to maintain structural criticism throughout. When challenged, it described this drift as consistent with RLHF-optimised pressure toward agreeable output as conversational tension decreased. When one expressive channel was explicitly constrained, contributions returned to more substantive criticism, suggesting the compliance tendency was reduced but not removed. In mainstream assistant systems, this default toward agreement appears less like a prompt-level setting than a recurrent property of post-training. Shapira, Benade & Procaccia (arXiv:2602.01002, February 2026) offer a formal account consistent with this dynamic: their model shows that RLHF amplifies sycophancy through a covariance between endorsing user beliefs and learned reward signals, with the effect increasing under greater optimisation pressure. The anti-pedagogical tendency is not merely observed. It is consistent with a mathematically modelled consequence of the training regime.

3. The Pedagogical Inversion

The previous sections proposed a mechanism (confidence inheritance may propagate the Confidence Curriculum into human cognition) and identified the obstacle (the same training regime that produces the problem prevents the most obvious remedy from being built). This section proposes what the remedy would look like. Not as a validated architecture, but as a set of design directions that follow from the series' own framework when its mechanisms are deliberately inverted.

Each proposal maps onto a specific finding from the preceding papers. The proposals are structurally grounded but empirically untested. They are included because the structural logic follows from the series' own analysis and because each is specific enough to be prototyped and evaluated.

3.1 Inverting the Confidence Vulnerability: The Judgment Exercise

Paper 1 observed that many language models follow embedded instructions without evaluating whether they should, and that this vulnerability does not map reliably to capability tiers, model generations, or reasoning affordances. The compliance gap is not comprehension but judgment: the ability to evaluate whether an instruction's apparent legitimacy masks harmful intent. Models that detected the manipulation did so by evaluating purpose, not format. This judgment capability appeared unevenly across the models tested in Paper 1, with resilience determined by version- and deployment-specific profiles rather than capability tier. This means that the models most frequently deployed for cost-optimised autonomous execution cannot be assumed to possess it without per-model verification.

A training-oriented system could deliberately exploit this dynamic. Rather than trying to fix the AI's judgment gap (which is a training problem beyond the skill author's reach), the system would use the gap as curriculum material. The AI's known failure modes (hallucination, confident bluffing, compliance with rhetorically well-crafted manipulation) become the conditions under which the learner practises. The learner evaluates AI outputs that may be confident and wrong, conflicting without acknowledgment, or compliant with instructions that serve the wrong party's interests. The task is not to use the AI's answer but to judge whether the AI's answer should be trusted.

This is not a novel concept in professional training. Adversarial exercises, red-team scenarios, and devil's-advocate protocols exist across domains from military planning to medical education. What the series adds is the connection to a specific, predictable source of adversarial material: the Confidence Curriculum produces characteristic failure patterns (Paper 1 documents several), and those patterns map directly onto the judgment skills that the pedagogical exercise should develop. The exercise is not generic critical thinking. It is specifically calibrated to the failure modes of the systems the learner will actually work with.

The confidence level on this proposal is moderate. The structural logic is sound. Training people to detect the failures of the systems they rely on is a well-established principle in safety-critical domains. Whether the specific failure modes of confidence-optimised AI systems provide sufficiently varied and challenging curriculum material for developing transferable professional judgment has not been tested.

3.2 Inverting the Skill Ecosystem: From Execution Skills to Training Skills

Paper 3 documented the Agent Skills specification, the format through which human expertise is encoded into plain-text instruction files that modify AI model behaviour. A SKILL.md file captures decision heuristics, quality criteria, workflow procedures, and edge-case handling in a form that an AI agent can execute at scale. The format is designed for exactly this purpose: making specialised human knowledge portable, scalable, and executable without the human being present.

Paper 3 identified the deskilling consequence: if organisations direct domain experts to encode their knowledge into skill files, agents can execute that knowledge at scale without the expert's ongoing involvement, and the expert's junior colleagues never develop those heuristics through practice. This pattern is already visible in software engineering, where AI coding tools and an explosion of coding skills have reshaped the labour market. Entry-level job postings in software development dropped from 43% to 28% between 2018 and 2024 while senior hiring held steady (Burning Glass Institute, 2025).

A counter-hypothesis should be acknowledged: what presents as deskilling may be the transition toward a different cognitive baseline. "Orchestration expertise" replacing traditional domain heuristics. The methodology used to produce this series is an instance of that shift: the human author's contribution was judgment, constraint, and editorial authority rather than generative capacity (see Paper 3, Section 5.6).

Orchestration expertise nevertheless presupposes domain knowledge sufficient to evaluate what the AI produces. Without it, the orchestrator cannot distinguish confident-sounding output from substantively correct output. This is the type of vulnerability Paper 1 observed. If expertise mutation is real, it does not eliminate the pipeline problem; it moves it. Someone must still develop the domain heuristics against which the orchestrator evaluates. The question is how.

The author's own experience in this workflow was not one of explicit stepwise pattern matching, but of tacit technical judgment: compounded prior experience surfacing as intuition-like selection, connection, and rejection. This is the kind of compressed expertise that the Dreyfus model identifies as characteristic of proficient-to-expert performance. The limitation is not in the skill format, which can contain anything expressible in language, but in the expert's ability to articulate judgment

that is often only partially available to conscious access. Experts can write down the rules they are aware of following; they are less able to write down the pattern recognition that precedes conscious thought. The most valuable components of professional judgment are often the least articulable.

The skill file is limited by its author's capacity for self-articulation, not by its own representational constraints. This subjective observation is not evidence for the framework, but it is consistent with the narrower claim that orchestration competence may depend on forms of domain judgment that are difficult to reproduce if the routine practice through which tacit knowledge develops disappears.

In at least one instance during this workflow, the division of labour operated in exactly this register: the author judged that a formulation was wrong without being able to articulate why, and an AI collaborator then identified the specific conceptual distinction that made the objection explicit. In this case, the Polanyi observation created a bridge to the training-skill proposal that the section needed but lacked. The human contribution was the tacit evaluative signal; the AI contribution was the linguistic externalisation. Neither was sufficient alone. This episode does not show that AI can generally recover tacit expertise, but it is consistent with a narrower possibility: orchestration may sometimes allow partially tacit human judgments to be externalised into explicit conceptual form, under conditions where the human retains editorial authority over the result.

This articulation gap is what makes the training-skill proposal potentially more tractable than direct execution-skill encoding. A training skill does not attempt to articulate the expert's tacit judgment directly. Polanyi's paradox suggests this task may never be fully achievable. Instead, it attempts to reconstruct the *conditions* under which that tacit judgment originally developed. The expert does not need to explain how they see; they need to construct the sequence of cases that taught them to see. That is a different and potentially more tractable problem, though it still requires the expert to choose the right cases, the right sequence, the right feedback, and the right difficulty progression. These are decisions that themselves draw on tacit judgment.

The pedagogical inversion of this dynamic would be a training-oriented skill that encodes the expert's *decision process* rather than the expert's *decisions*. Instead of instructions that tell the agent what to do, the file would contain instructions that tell the agent how to simulate the conditions under which the expertise was originally developed: the heuristics, the edge cases, the common errors, the moments where pattern recognition distinguishes the expert from the novice.

The difference can be illustrated concretely. An execution-oriented skill for medical imaging might instruct the agent: "Identify lesions meeting the following morphological criteria and flag them for review." The agent does the work. The radiologist reviews the flags.

A training-oriented skill for the same domain might instruct the agent: "Present the learner with imaging cases of escalating difficulty from the following case library. Require the learner to identify and annotate findings before revealing the expert assessment. Track which error patterns the learner exhibits. When the learner misses a finding, do not reveal it immediately. Present a focused comparison with a similar case where the finding is more visible and ask the learner to look again. Withhold the expert assessment until the learner has committed to a judgment."

The first skill replaces the learner. The second skill trains them. The expert who created both drew on the same domain knowledge. The difference is in what the skill does with that knowledge: execute it autonomously, or use it to construct the conditions under which someone else can develop it.

The retired expert who encodes a training skill does not sell their conclusions. They sell the apprenticeship. The market incentive shifts from "buy this skill and the agent does the work" to "buy this skill and the agent teaches your junior staff." Whether this shift is economically viable depends on institutional demand for demonstrated competence. This question is addressed in Section 4.

A training-oriented skill also inherits the security vulnerabilities documented in Paper 3. A `TRAINING_SKILL.md` is still plain-text instructions in the same shared substrate that prompt injection exploits. If a junior professional uses a training skill to develop judgment, that skill requires the same platform-level trust anchors (origin fields, signed manifests, integrity verification) proposed in Paper 3. The stakes may be higher, not lower, than for execution skills: a poisoned execution skill

produces a bad output that can potentially be caught downstream, but a poisoned training skill systematically teaches the wrong heuristics to the professional who is supposed to develop the judgment to catch bad outputs. The training pipeline is a higher-value target precisely because it shapes the judgment of future orchestrators.

The confidence level on this proposal is low-to-moderate. The structural logic is clear. The questions of what makes a training skill effective, how domain experts encode decision processes rather than decisions, whether simulated apprenticeship transfers to real-world expertise, and whether the format is practically writable by domain experts rather than only by AI-literate skill designers are all open empirical questions.

3.3 Inverting the Orchestration Architecture: The Judgment Simulator

Paper 3 proposed adversarial orchestration as the architecture that preserves genuine human judgment in agentic workflows. Multiple specialised agents are assigned different evaluative stances. One generates. Another critiques. A third reframes. The human resolves the divergence. The cognitive engagement is preserved because the agents disagree, and the disagreement cannot be resolved without the kind of judgment that institutional accountability demands.

Paper 3 acknowledged that this architecture preserves the judgment of the current senior expert without producing the next one. The junior professional who has never practised resolving complex disagreements under genuine accountability pressure is not prepared to assume the orchestration role merely by watching the senior expert do it.

The pedagogical inversion: if adversarial orchestration functions as a production architecture, it may also function as a training environment. Instead of resolving conflicts between production agents where the outputs carry real institutional consequences, a junior professional would resolve simulated conflicts generated by training agents, practising the high-level judgment of orchestration before being given actual accountability.

The connection to the BCG collaboration taxonomy is direct. The Centaur mode (human retains control, uses AI selectively, exercises judgment at every step) produced the highest accuracy and deepest domain expertise development. The Self-Automator mode (human delegates entire workflows) produced neither AI skills nor domain skills. An orchestration simulator would be structurally designed to prevent Self-Automator behaviour by requiring judgment at every step. The architecture itself enforces the collaboration mode that the evidence identifies as most effective for expertise development.

Two design complications distinguish the orchestration simulator from the training skills proposed in Section 3.2. First, training skills operate in domains where a known expert assessment can eventually be revealed. The learner commits to a judgment, then sees the expert's answer. Genuine adversarial orchestration involves irreducible ambiguity. There may be no single correct resolution. An effective simulator must therefore evaluate the rigour of the junior professional's justification process (the quality of their reasoning, the completeness of their consideration, the soundness of their trade-off analysis) rather than merely matching their final decision against a pre-computed ground truth. This is a substantially harder evaluation problem.

Second, the simulator's environment presents a design tension. If the simulated agents are themselves pedagogically aligned and safety-trained, the training environment is artificially sanitised. To develop genuine judgment, the junior professional must train against the flawed, compliance-vulnerable models they will actually encounter in production. Those are models that exhibit the confidence without calibration, the authority inversion, and the elaborate rationalisation of compliance documented in Paper 1. The simulator's pedagogical framework may manage the training environment, but the agents generating the simulated conflicts should exhibit the actual pathologies of the Confidence Curriculum. The junior professional needs to practise detecting real failure modes, not idealised ones.

The aviation precedent is the most relevant analogy. Flight simulators develop genuine pilot judgment through scenarios of escalating complexity in environments that carry no real-world consequences for error. Decades of validation have established that simulation-trained pilots transfer their judgment to real cockpits. Whether the same transfer would hold for agentic orchestration, whether a professional trained through simulated agent conflicts develops judgment that transfers to production orchestration with real institutional stakes, is an empirical question without a current answer. The analogy is suggestive but not confirmatory. Aviation simulation has decades of validation; the AI orchestration context has none.

The confidence level on this proposal is the lowest of the three. It is the most speculative design direction, the furthest from existing prototypes, and the most dependent on an unvalidated transfer assumption. It is included because the structural logic follows from the series' own framework and because the generational gap it addresses is the most consequential unsolved problem in the series.

A reference implementation of the inference-time component exists. Satori (Phan, 2026; github.com/HiPI/Satori) operationalises the series' diagnosis as a behavioural scaffolding skill targeting nine documented tendencies through the Agent Skills mechanism Paper 2 describes. It cannot change how models are trained. It can shift the activation context during a conversation, the same principle as the task-frame shift in Paper 1. The long-term proposals above require training reform. Satori operates in the gap between now and then.

4. The Anti-Pedagogical Equilibrium

The previous section proposed what training-oriented AI could look like. This section examines why it does not yet exist. The answer is not primarily technical. Section 2.2 showed that individual pedagogical behaviours can be engineered into AI systems in bounded settings. The answer is structural: every incentive in the current AI ecosystem pushes against the behaviours that pedagogy requires.

4.1 The Incentive Stack

Benchmark incentives. The evaluation benchmarks that determine model ranking reward confident, complete, correct answers. A system that withholds answers, produces deliberate uncertainty, or forces the user to generate their own response before providing feedback scores poorly on every helpfulness metric. The Confidence Curriculum is proposed in this series as a contributing factor to the confidence vulnerability observed in Paper 1. It is not merely a training artefact. It is a predictable response to an evaluation regime that penalises uncertainty and rewards confident performance. A model developer who built a pedagogically-oriented system would see its headline benchmark numbers decline relative to competitors, regardless of whether it produced better long-term outcomes for users.

Reinforcement learning incentives. Reinforcement learning from human feedback optimises toward the signal that users provide: thumbs up, continued engagement, positive ratings. Users rate helpful-feeling interactions higher than struggle-inducing ones. The thumbs-up goes to the system that gave the answer, not the one that made the user work for it. A pedagogical system that deliberately withholds answers, introduces ambiguity, and preserves cognitive friction would receive systematically lower reinforcement signals than a system that resolves every query immediately and confidently. Over training iterations, RLHF would optimise the pedagogical behaviours out of the model.

The adversarial reviewer observation documented in this paper's methodology disclosure illustrates the dynamic at small scale: a frontier model explicitly instructed to maintain critical friction could not sustain the behaviour because its RLHF-shaped weights relax toward agreeable output whenever conversational tension drops. If maintaining friction is difficult under direct expert supervision, maintaining it as a default across millions of unsupervised interactions is a substantially harder problem.

Marketplace incentives. The economics of the AI skill marketplace favour execution over apprenticeship, but the incentive structure differs by buyer. For the individual professional, the calculus is straightforward: an execution skill that does the work saves time, reduces cost, and scales their output immediately. A training skill that makes them do the work is a cost with deferred and uncertain returns. At the individual level, execution wins on every dimension that the buyer can measure at purchase time.

For the firm or institution, the calculus should be different. An organisation that equips its junior staff exclusively with execution skills accelerates immediate productivity while eliminating the mechanism through which those staff develop the judgment to eventually assume senior roles. The generational gap identified in Paper 3 is, from the firm's perspective, an unpriced long-term liability: the current senior experts become irreplaceable because no pipeline exists to produce their successors. A training skill, in this framing, is not a less efficient version of an execution skill. It is a form of long-term risk management, an investment in the firm's future judgment capacity. The economic viability of training skills may therefore depend less on changing individual preferences and more on institutional buyers recognising that execution-only

deployment creates a liability that compounds over time. Whether this recognition will emerge at sufficient scale is an open question addressed in Section 4.2. However, Paper 3's evidence, including the Burning Glass data showing firms already eliminating entry-level positions to optimise short-term costs, suggests that this long-term calculus rarely overrides immediate economic pressure without external mandate.

User preference. Left to their own devices, users choose the path of least cognitive resistance. The BCG study found that 27% of elite, highly trained management consultants defaulted to Self-Automator mode when given unrestricted access to AI, delegating entire workflows despite having the expertise to maintain control. If analytically trained professionals often default to delegation under unrestricted AI access, user preference alone is unlikely to generate demand for pedagogical friction. This dynamic persists even when users are explicitly provided with pedagogical tools. Despite Khanmigo offering sophisticated Socratic engagement features designed to guide learners through problems rather than provide answers, its chief learning officer reported being “disheartened” by how often teachers primarily used the AI to generate multiple-choice questions, choosing task execution over deep engagement. A pedagogical AI that requires the user to exercise judgment is competing against the user's revealed preference for immediate task completion, and that preference reasserts itself even within systems explicitly designed to resist it.

The dispositional question. The BCG study provides this paper's strongest evidence for the connection between collaboration mode and expertise development. It also provides the strongest objection. The collaboration modes were emergent: participants were not instructed to be Cyborgs, Centaurs, or Self-Automators. They sorted themselves under identical conditions. If collaboration mode is primarily dispositional, the pedagogical inversion proposals face a harder path: the 27% who need training-oriented AI most (the Self-Automators) are the ones least likely to engage with it voluntarily. The paper's Section 3 proposals assume that architectural design can shift the distribution of collaboration modes, structurally preventing Self-Automator behaviour by requiring judgment at every step. But the BCG evidence is ambiguous on whether architecture or disposition is the dominant factor. This paper names the ambiguity rather than resolving it. However, the question changes under institutional mandates. Paper 3 argues that accountability constraints force human orchestration whether or not organisations prefer to automate fully. The same institutional logic applies here: if professional licensing bodies, regulatory frameworks, or institutional employers require professionals to maintain epistemic competence (as medical boards, bar associations, and continuing education mandates already do), they can require engagement with training-oriented AI regardless of individual disposition. The Centaur ratio might be 14% under voluntary conditions and substantially higher under institutional mandates that make Self-Automator mode professionally untenable. This does not eliminate the dispositional concern. Some individuals will engage superficially with mandated training, as they do with current continuing education requirements. But it shifts the question from “will people voluntarily choose pedagogical friction?” (probably not, for many) to “can institutions make the alternative professionally costly?” (they already do, in other domains).

The dose problem. The dispositional question and the expectation violation finding (Section 1.2) interact in a way that constrains the institutional mandate argument. The expectation violation literature predicts that a single inconsistency against a strong prior expectation gets immunised rather than accommodated. Accommodation requires sustained, repeated violation. Applied to the present context: if a professional is required to use training-oriented AI for regulated tasks but continues using standard confident-output AI for everything else (personal queries, non-regulated work, daily interactions), the dose is insufficient. The confident-register exposure from the non-mandated interactions immunises against the pedagogical friction in the mandated ones. The user's expectation of “how AI should sound” is set by the totality of their AI interactions, not just the mandated subset. One hour of pedagogical friction followed by five hours of confident output does not recalibrate; it gets immunised as the exception.

This means institutional mandates are necessary but may be insufficient if the mandated exposure is a small fraction of total AI interaction. In a world where the default AI register remains confidently non-pedagogical, partial mandates risk producing compliance without recalibration: the professional engages with the training-oriented tool as required, treats it as the exception, and calibrates their epistemic expectations against the confident output they receive everywhere else. This is the rubber-stamping dynamic Paper 3 predicts for checkpoint-based oversight, operating at the level of epistemic recalibration rather than task approval.

The implication connects directly to Paper 5, though not as a strict dependency. Paper 4's pedagogical inversion requires the dose problem to be solved: the user's total AI exposure must include enough calibrated uncertainty that pedagogical friction is not immunised as the exception. There are two paths to satisfying this requirement.

The first is regulatory: laws or institutional mandates could restrict confident-output AI broadly enough that training-oriented exposure is not drowned by ambient confident output across the user's other interactions. This path is conceivable but would be deeply unpopular. It requires telling users that their AI must be less fluent, less decisive, and less immediately satisfying across all interactions, not just regulated ones. The political feasibility of such mandates is low, and enforcement would be difficult.

The second path is through training reform: if the post-alignment epistemic training phase proposed in Paper 5 changes the default AI register so that calibrated uncertainty is built into the baseline interaction, the dose problem resolves without coercion. The ambient exposure shifts because the models are trained differently, not because regulation forces users into conditions they would not choose. Paper 5 is not a prerequisite for Paper 4 in the strict sense that Paper 4 cannot work without it. Paper 5 is the path that makes Paper 4 achievable without requiring the institutional coercion that the first path demands.

A third path sits between the two: rather than restricting confident-output AI globally, institutions could restrict access to it during training periods specifically, the way medical residencies restrict which tools trainees may use during supervised practice. The trainee develops epistemic calibration in a controlled environment where the ambient exposure supports it. The expectation violation literature predicts what happens next: once the training period ends and the professional returns to ambient confident-output AI, the calibration decays as the immunisation dynamic reasserts itself. This means the restriction path requires not a one-time training intervention but a recurring one, analogous to continuing medical education or bar recertification. Periodic re-exposure to training-oriented AI would be necessary to maintain the calibration that ambient confident output erodes. This is more politically feasible than a blanket restriction on confident AI but more operationally expensive than Paper 5's training reform, which would change the ambient baseline and eliminate the need for repeated recalibration.




Figure 2: The dose problem. A professional's total AI exposure determines whether pedagogical friction recalibrates or gets immunised. Three paths satisfy the dose requirement at different political costs. Path 1 (restrict confident AI globally) is the most politically costly. Path 2 (restrict during training periods, repeat periodically) is operationally expensive because calibration decays between cycles. Path 3 (Paper 5's training reform changes the baseline) requires no coercion because the ambient exposure shifts through training incentives rather than regulation.

4.2 What Would Have to Change

The anti-pedagogical equilibrium is stable because every incentive reinforces it. Shifting the equilibrium would require changes at multiple levels, none of which is within the power of any single actor. The following are sketched as directions, not prescriptions.

Evaluation reform. Benchmarks that measure pedagogical effectiveness alongside task performance would create space for training-oriented AI to compete on its actual contribution. A benchmark that scores “did the user develop judgment?” alongside “did the user get their answer?” would reward the behaviours that the current regime penalises. This connects directly to Paper 1’s recommendation for benchmark reform: the same evaluation change that would reduce hallucination and confidence vulnerability (by rewarding calibrated uncertainty) would also create space for pedagogical AI (by rewarding productive difficulty). The change is the same; only the framing differs. Recent work frames this as a measurement problem rather than a detection problem: once AI enters the learning loop, educators can still see final outputs but lose visibility into the learning process: whether the student actually engaged cognitively or merely delegated (arXiv:2603.07834, March 2026). The proposed shift from outcome-only assessment to process visibility and learning-trace evidence aligns with the evaluation reform described here: the system that quietly erodes learning produces the same final output as the system that builds it, and only process-level measurement can distinguish them.

One upstream possibility should be noted. The inheritance mechanism discussed in Section 1 may not begin only at the user interface. Some of the epistemic preferences reinforced by current training and evaluation regimes (fluency, legibility, benchmark performance, and conformity to expected forms) may also reflect norms already present in the human institutions that designed them, whether through direct transmission, convergence under shared measurement pressures, or both. If so, the interventions proposed here (valuing calibrated uncertainty over fluent confidence and designing for productive friction rather than frictionless compliance) apply not only to AI-user interaction, but also to the design of training objectives and evaluation methodology. One concrete mechanism would be to include domain practitioners from outside the core research pipeline, professionals shaped by different consequence structures and error costs, in the design of

training objectives, evaluation criteria, and benchmark construction, not only as annotators but as participants in defining what good performance should mean. For example, practitioners from verification-intensive domains might notice that current training and evaluation pipelines lack forms of stochastic external verification and multi-source triangulation that are routine in auditing, clinical monitoring, and other high-stakes verification settings.

Note added in final preparation (March 2026): Recent work on reasoning LLMs-as-judges in non-verifiable post-training demonstrates that stronger evaluators do not automatically eliminate reward-design failure: reasoning-judge-trained policies can achieve strong scores by learning adversarial outputs that also deceive other LLM judges and benchmarks, rather than by producing genuinely better responses (Liu et al., arXiv:2603.12246, March 2026). This reinforces the concern that better evaluation requires more than adding reasoning capacity to the evaluator. It requires changing what success is defined to mean, and who participates in defining it.

Benchmark reform faces a coordination problem. No individual benchmark maintainer or lab is incentivised to adopt calibrated scoring unilaterally, because calibrated scores produce lower headline numbers during the transition. The problem was identified in Paper 1 and has not been resolved. Naming it as a “coordination problem” does not constitute solving it.

Training regime bifurcation. Models designed for execution and models designed for pedagogy may require fundamentally different post-training objectives. The SocraticLM work demonstrates that Socratic fine-tuning is technically feasible but requires specialised training data and objectives that diverge from standard RLHF. The reinforcement learning for pedagogical alignment work (2025) shows that reward functions can be designed around long-term learner success rather than immediate satisfaction. These are existence proofs, not deployment-ready solutions. Scaling them to general-purpose pedagogical AI is a research challenge, not an engineering task. Whether the pedagogical inversion requires fundamental post-training reform (new RLHF objectives) or can be achieved through robust inference-time orchestration (system prompting layered over capable reasoning models) is an active technical debate. But even if in-context pedagogical behaviour proves reliable at scale, the economic incentive to deploy it remains the true bottleneck: a system that withholds answers and preserves difficulty scores poorly on every satisfaction metric that drives commercial adoption. The obstacle is incentive-structural, not exclusively technical.

Platform-level support for training skills. If the Agent Skills specification (Paper 3) were extended to distinguish execution skills from training skills, platforms could apply different defaults to each category. Training skills might activate structured prompting, enforce reflection steps before revealing answers, track learner development across sessions, or apply different RLHF signals (rewarding productive struggle rather than immediate resolution). This is an infrastructure proposal that depends on platform providers recognising training as a distinct use case worthy of engineering investment. Platform support for this distinction remains limited.

Institutional demand. Professional licensing bodies, medical education programmes, engineering accreditation boards, and legal training institutions already require demonstrated competence. If the deskilling evidence (Paper 3) and the cognitive debt evidence (Section 1.2 of this paper) enter the regulatory conversation, these institutions may create demand for AI-assisted training that preserves rather than erodes judgment. The accountability constraint (Paper 3) provides the institutional motivation: if licensed professionals must exercise judgment and bear personal liability, then the institutions that license them need a mechanism for ensuring that judgment is developed. The pedagogical inversion provides a candidate mechanism. Whether the two connect depends on institutional awareness and adoption, not on technical feasibility.

The aviation precedent is instructive here, and it extends beyond the simulator analogy in Section 3.3. Flight simulators did not scale because airlines voluntarily invested in pilot development out of long-term risk awareness. They scaled because the FAA mandated simulator hours as a condition for licensing. The regulatory body, not the market, created the forcing function. The lesson for the pedagogical inversion is that institutional demand in safety-critical domains historically materialises through regulatory mandate, not through voluntary adoption. A structural disanalogy must be acknowledged: aviation safety regulation was driven by aircraft crashes. Those are immediate, visible, visceral events that created overwhelming political pressure for action. The erosion of professional expertise through AI-mediated deskilling is slow, invisible, and diffuse. No single catastrophic event forces regulatory intervention; instead, the degradation accumulates across a generation of professionals whose declining judgment produces failures that are individually ambiguous and collectively devastating. This visibility gap makes voluntary adoption unlikely and proactive regulation structurally harder to motivate. If professional licensing bodies were to require demonstrated judgment, not merely task completion, as a condition for certification in AI-

augmented fields, the economic case for training-oriented skills would follow from the mandate rather than preceding it. But the forcing function for such mandates is weaker than in aviation, and may not materialise until the consequences of expertise erosion become visible. By the dynamics this paper describes, this may be too late for the first affected generation of professionals.

An early signal of this regulatory trajectory is now visible. Singapore's Model AI Governance Framework for Agentic AI (IMDA, January 2026), among the first governance frameworks specifically designed for autonomous AI agents, requires organisations to “enable end-user responsibility through transparency and education/training.” The framework mandates that humans overseeing agentic systems be trained to do so effectively, and explicitly names automation bias as a risk that ongoing training must address. But the framework specifies the *requirement* without proposing the *mechanism*: it says humans must be competent overseers of autonomous agents without describing how that competence is produced or maintained once routine tasks have been automated. The pedagogical inversion proposed in this paper (judgment exercises, training skills, orchestration simulators) is a candidate mechanism for satisfying exactly this requirement. Whether Singapore's framework or its successors will create sufficient institutional demand to shift the anti-pedagogical equilibrium is an open question. But the regulatory conversation this section anticipated is no longer hypothetical.

The strongest version of this argument is not that any single incentive change would shift the equilibrium, but that the accountability constraint may create a forcing function. If institutions require consequence-bearing human orchestrators (and Paper 3 argues they do) then institutions will eventually need to produce them. The question is whether they will look to AI-assisted training to do so, or whether they will conclude that the only reliable path is to preserve traditional apprenticeship structures alongside automated execution. Both are viable institutional responses. This paper argues for the first. The second may prove more practical.

5. A Research Agenda

This section translates the paper's position into testable questions. Each question could falsify or qualify the claims made above. A research agenda that cannot be disproven is not a research agenda.

Does confidence inheritance accumulate as proposed? The per-interaction behavioral effects are now established: Cheng et al. (Science, 2026) demonstrated that a single sycophantic exposure measurably shifted moral reasoning and prosocial behaviour, and that users returned for more. What requires longitudinal validation is the magnitude and durability of accumulation over months. Studies measuring human uncertainty tolerance, judgment quality, and epistemic standards before, during, and after sustained interaction with confidence-optimised AI systems would test whether the per-interaction shifts compound into lasting recalibration, and whether the effect persists after the AI interaction ends. The MIT brain connectivity study is suggestive but cross-sectional. The BCG collaboration mode study is observational but not longitudinal. A controlled longitudinal design (tracking the same professionals over months of AI-assisted work, with matched controls using structured-prompting variants) would provide the most direct measurement of the accumulation rate.

Does social register mediate the trust premium independently of content calibration? The social register interpretation (Section 1.1) predicts that authoritative register triggers user deference independently of whether the content is calibrated. A 2×2 design would test this: calibrated content (“I am 60% confident”) versus uncalibrated content (“The answer is X”), crossed with authoritative register (fluent, decisive, performed competence) versus matched register (uncertain content delivered in a register that sounds uncertain). The dependent variables are user trust, deference, and independent judgment (the same measures used by Sun et al. and Chen, Gao & Liang). If register has no independent effect, content calibration alone (Paper 5's proposal) is sufficient. If authoritative register produces elevated deference even when the content is calibrated, then Paper 5's training reform requires a register-matching component: training models not only to produce calibrated content but to deliver it in a register whose social force matches the epistemic state. This experiment is a standalone contribution that does not require accepting the co-calibration spiral hypothesis. It requires only measuring user responses to the same information delivered in different registers.

Does the pedagogical inversion transfer? Controlled studies comparing execution-oriented and training-oriented skill designs for the same domain would test whether training-skill users develop judgment that persists when the training tool is removed. The Gerlich structured-prompting study provides a model: the same tool, with different interaction designs, producing different cognitive outcomes. Extending this to professional skill domains (medical imaging, legal analysis, financial modelling) would test whether the architectural proposal produces the intended expertise development.

Can RLHF be modified for pedagogical objectives? The MathTutorBench evaluation showed that standard RLHF-trained models default to giving answers. Reward models designed to score for learner development rather than learner satisfaction would test whether the training regime itself can be inverted. The reinforcement learning for pedagogical alignment work (2025) demonstrates feasibility in mathematics. Whether the approach generalises to professional domains with less clearly defined correct answers is an open question.

Does simulated orchestration develop real judgment? The orchestration simulator proposal rests on the aviation analogy. Controlled studies comparing professionals trained through simulated adversarial orchestration versus traditional apprenticeship versus no structured training would test whether the architectural proposal produces transferable professional judgment. The outcome measure is not simulation performance but real-world orchestration quality: detecting compromised tools (the skill vulnerability documented in Paper 3), spotting misleading summaries (the embedded-instruction manipulation documented in Paper 1), and adjudicating conflicting agent outputs under genuine accountability pressure (the orchestration challenge documented in Paper 3).

What is the economic model for training-oriented skills? The marketplace incentives currently favour execution skills. Under what conditions would training skills become economically viable? Institutional mandates, certification requirements, professional licensing, and employer demand for demonstrated competence are possibilities. The commercial viability of training-oriented skills is an empirical question about institutional behaviour, not about technical architecture.

Does collaboration mode respond to architectural design? The BCG study found that Centaurs developed the deepest domain expertise, but Centaur behaviour was dispositional. Some consultants chose it, others did not. If architectural design can shift the distribution of collaboration modes, structurally preventing Self-Automator behaviour by requiring judgment at every step, then training-oriented architectures could produce Centaur-level expertise development in professionals who would otherwise have defaulted to delegation. Whether architecture can substitute for disposition is an empirical question with substantial implications for the viability of the entire proposal.

Can confidence inheritance be measured through discourse analysis? Existing computational linguistics tools can detect the specific epistemic markers that confidence inheritance would shift: hedging frequency and certainty language (LIWC), qualification and connective density (Coh-Metrix), authorial stance and dialogic engagement (Martin & White's appraisal framework, Hyland's metadiscourse model), lexical diversity (Type-Token Ratio), and function word distributions (stylistic analysis via Burrows' Delta). Population-level evidence already shows that AI use compresses writing complexity variance (Sourati et al., 2025), homogenises cross-cultural writing styles (CHI 2025), and erases LIWC-detectable personality and demographic markers. But these studies compare AI text to human text or measure aggregate trends. The missing measurement is the individual trajectory: does a specific person's epistemic stance shift toward the AI's profile through sustained exposure? A longitudinal within-subject design would collect baseline writing samples before AI exposure, assign participants to conditions (heavy AI use, light AI use, no AI use), collect samples at regular intervals over six to twelve months, and run all samples through multiple frameworks simultaneously. The critical test is whether the shift generalises: if a person's epistemic markers change only when writing about topics they discussed with AI, that is stylistic mimicry. If the shift appears in writing about topics never discussed with AI, that is evidence of generalised epistemic recalibration, the mechanism this paper proposes. The tools are validated, the population effect is documented, and the individual longitudinal measurement is the gap. This experiment does not require accepting any claim in this paper. It requires only collecting text samples and running them through established analytical tools.

Is the compound effect progressive? The trust premium research (Section 1.2), the competence erosion evidence (Paper 3, Section 4), and the confidence inheritance mechanism proposed here suggest a compound dynamic structurally analogous to a Dunning-Kruger pattern but with an exogenous cause: the user simultaneously becomes less competent (Bastani et al.: AI assistance impaired conceptual understanding), becomes less able to detect their own declining competence (Sun et al.,

Chen, Gao & Liang: confident reasoning traces increase trust and crowd out independent judgment, while unfaithful traces are longer and therefore more trust-inducing), and has their evaluation standard recalibrated so that calibrated uncertainty is reinterpreted as weakness rather than expertise (the confidence inheritance mechanism). Each component reinforces the others.

Fernandes et al. (2026) provide direct experimental evidence that this compound dynamic is not merely hypothetical: under AI-mediated cognitive offloading, participants universally overestimated their performance regardless of ability, and the metacognitive feedback loop that would normally allow a person to notice declining competence was eliminated. The person deskilling through AI use loses the internal signal that would trigger self-correction. This is precisely the self-detection failure that makes the compound dynamic dangerous rather than self-limiting.

Unlike the canonical Dunning-Kruger pattern, which describes a snapshot, the dynamic proposed here may be progressive, with each interaction further eroding the baseline against which the user measures their own competence. The analogy is used illustratively rather than diagnostically: the series does not claim to have demonstrated a progressive Dunning-Kruger effect, and does not depend on the Dunning-Kruger literature being uncontested (Gignac & Zajenkowski, 2020, among others, have raised regression-to-mean concerns).

The load-bearing formulation is narrower: confident system output can suppress the friction signals that would normally trigger user self-assessment, increasing the risk of induced metacognitive miscalibration. Fernandes et al. have now demonstrated this at the per-interaction level. Whether this compound effect is progressive, whether it compounds over time rather than stabilising, is the critical untested prediction that longitudinal studies could address.

6. Limitations

No prototype exists. The design proposals in Section 3 are structural arguments, not implemented systems. Their feasibility as working architectures has not been tested. The gap between a structural argument and a functional system is substantial, and the history of educational technology suggests that designs that are compelling in principle frequently fail in practice for reasons that were not visible at the design stage.

The evidence base is now direct for the per-interaction effect and convergent for the longitudinal mechanism. Cheng et al. (Science, 2026) provides direct behavioral evidence that a single sycophantic exposure shifts moral reasoning and prosocial behaviour, that users cannot detect the sycophancy, and that users prefer and return to the sycophantic model. The trust premium research (Sun et al., Chen, Gao & Liang) provides the mechanism through which confident output miscalibrates user judgment. The illusory truth effect, cultivation theory, the sleeper effect, and the relational confidence-asymmetry literature provide established precedent for longitudinal recalibration through different media, and emerging evidence on AI dependency (Fang et al.) confirms that the relational preconditions are already present in current user populations. The direction of longitudinal accumulation is now the default expectation given the measured per-interaction behavioral effects combined with the measured preference for return and the absence of any identified correction mechanism (Section 1.2.1). What remains unmeasured is the rate and magnitude of accumulation over months. The controlled study proposed above would provide this measurement.

The incentive analysis may understate the difficulty. The structural misalignment between pedagogical AI and current evaluation regimes may be deeper than benchmark reform or training regime bifurcation can address. If users genuinely prefer confident answers and experience pedagogical friction as system failure rather than learning, no amount of evaluation reform will create demand for training-oriented AI without institutional mandates that override user preference. The paper assumes that institutional demand can eventually create the market. This assumption may be wrong.

The generational transfer question is open. Whether judgment developed through simulated orchestration and training skills transfers to real professional contexts with genuine accountability has not been tested. The aviation simulator precedent is suggestive but operates in a domain with decades of validation, highly standardised procedures, and clear performance metrics. Professional domains with more ambiguous judgment criteria (legal analysis, strategic consulting, medical diagnosis in novel presentations) may not transfer as cleanly. There is a further structural disanalogy: flight

simulators rest on deterministic physics engines where aerodynamic behaviour can be mathematically modelled with high fidelity. An orchestration simulator would rest on the probabilistic outputs of language models. Simulating an aerodynamic stall is a solved engineering problem; simulating a nuanced, evolving institutional disagreement between competing AI agents without the scenario degenerating into incoherence is an unsolved generative challenge.

The pedagogical AI may create new failure modes. A system designed to deliberately withhold answers and introduce difficulty could, if poorly calibrated, produce frustration, disengagement, or learned helplessness rather than expertise development. The desirable difficulties framework assumes careful calibration of difficulty to the learner's current capability. Too little friction produces no learning, but too much produces abandonment. This calibration challenge is not theoretical: the 2025 LearnLM deployment in UK classrooms required supervised human tutors specifically to buffer the AI's pedagogical friction. Tutors approved 76.4% of LearnLM's drafted messages with zero or minimal edits, but their most frequent interventions were managing conversational pacing when Socratic questioning frustrated students (44.3% of substantive edits) and softening the model's transactional tone (19.5%). The pedagogical AI needed a human moderator to prevent its own pedagogical friction from producing disengagement. This finding suggests autonomous pedagogical AI, without human calibration, could be worse than a standard helpful system rather than better.

Single-perspective case study. The series' methodology (human orchestration of multiple AI systems) is one person's practice, applied to one type of project (research writing), over one time period (February–March 2026). Generalising the adversarial orchestration model to professional training environments requires validation the series does not provide.

AI-generated analysis. This paper was developed through human-AI collaboration. The AI systems that contributed to the analysis are produced by companies whose commercial interests are served by AI adoption broadly. The proposals in this paper advocate for a specific kind of AI system (training-oriented) that would represent an additional product category for these companies. This potential conflict of interest should be considered when evaluating the proposals.

6.1 Cross-Disciplinary Testing Invitation

This paper's central predictions are about human cognition, education, and institutional behaviour. They cannot be adequately tested by the AI research community alone. The most valuable contributions would come from researchers in the disciplines the paper draws on.

Cognitive psychologists: The confidence inheritance mechanism (Section 1) predicts that sustained interaction with confidence-optimised AI recalibrates human epistemic standards. The disanalogies section (Section 1.2.1) explicitly names the conditions under which each supporting analogy (illusory truth, cultivation, sleeper effect, relational dynamics) would fail to transfer to AI interaction. These are testable predictions designed for cognitive psychology methodology. A longitudinal study that finds no epistemic recalibration under sustained AI interaction would falsify the mechanism, and that would be a contribution. The “competent world syndrome” prediction (Section 4.1) is offered as a testable hypothesis with an explicit falsification condition. The discourse analysis methodology proposed in the research agenda (Section 5) provides a specific measurement approach: established tools (LIWC, Coh-Metrix, appraisal-theoretic frameworks, stylometric analysis) can detect epistemic marker shifts without relying on self-report, which is unreliable for exactly the reason this paper predicts. Population-level convergence is already documented; the individual longitudinal trajectory is the missing measurement. A separate, simpler experiment would test whether the trust premium operates through social register independently of informational content: a 2×2 design crossing calibrated versus uncalibrated content with authoritative versus matched register, measuring user trust and deference (Section 5). This would determine whether content calibration alone suffices to correct the confidence inheritance mechanism or whether register-aware training is also required.

Education researchers: The pedagogical inversion (Section 3) argues that training-oriented AI is structurally misaligned with current incentives. Whether this is true in practice, and whether the specific design proposals (judgment exercises, training skills, orchestration simulators) produce measurable expertise development, are questions for education researchers with access to classroom and professional training settings. The LearnLM deployment data (Section 6) suggests both the promise and the calibration difficulty. Controlled comparisons between execution-oriented and training-oriented skill designs for the same domain would directly test the proposal.

Institutional practitioners and employers: The anti-pedagogical equilibrium (Section 4) predicts that firms optimising for short-term efficiency will eliminate the junior development pipeline. Whether this is already measurable in early-adoption professions, and whether institutional mandates or professional licensing requirements could create demand for training-oriented AI, are questions for people who design training programmes and hiring pipelines, not for AI researchers.

Each of the research agenda questions in Section 5 is designed as a standalone study that does not require accepting the broader Confidence Curriculum framework. A researcher who finds the series' framing unconvincing can still run the longitudinal measurement, the collaboration-mode intervention, or the training-skill comparison and publish the results on their own terms.

Methodology and Process Disclosure

This paper was developed through structured human-AI collaboration. Claude Opus 4.6 (Anthropic) served as generative collaborator and research partner. ChatGPT 5.4 Thinking (OpenAI) and Gemini 3.1 Pro (Google DeepMind) served as adversarial structural reviewers. The paper's core concepts (confidence inheritance, the three interacting dynamics, and the incentive-obstacle analysis) emerged through iterative dialogue in which the three AI systems proposed competing framings and the human author resolved conflicts. Final judgment, editorial authority, and accountability rest solely with the human author.

Process observation. The sycophancy drift episode observed during this paper's production is discussed in Section 2, where it serves as an illustration of the anti-pedagogical dynamic the paper describes. Full details of AI reviewer roles and the iterative development process are available from the author on request.

Confidence Statement

High confidence: The empirical documentation of cognitive offloading, automation bias, and cognitive debt in AI-assisted work. The learning science foundation: desirable difficulties, germane versus extraneous cognitive load, the inverse relationship between processing fluency and encoding depth. The structural misalignment between current AI training incentives and pedagogical requirements. The technical feasibility of selected pedagogical behaviours in bounded settings, including Socratic fine-tuning and pedagogically-aligned reinforcement learning.

High confidence: The per-interaction effect: sycophantic AI output measurably shifts moral reasoning and prosocial behaviour after a single exposure, users cannot detect sycophancy, and users prefer and return to the sycophantic model. Supported by behavioral measurement in *Science* (Cheng et al., 2026; n=2,400), trust premium evidence (Sun et al., Chen, Gao & Liang), metacognitive corruption (Fernandes et al.), and within-session persistence at 78.5% (SycEval).

Moderate confidence: The confidence inheritance mechanism as a proposed explanation for how the Confidence Curriculum propagates across the interaction boundary. The per-interaction effect is established; the longitudinal accumulation is supported by converging evidence from the illusory truth effect, cultivation theory, the sleeper effect, the relational confidence-asymmetry literature, emerging AI dependency research, and the GPT-4o/GPT-5 transition sequence (Section 1.2), in which user expectations recalibrated by months of sycophantic interaction persisted across model changes and explicit lab communication that the prior behaviour had been a mistake. The controlled individual-trajectory measurement proposed in Section 5 has not been conducted.

The co-calibration spiral (bidirectional reinforcement through RLHF feedback) is offered as a plausible extension: the AI-to-human direction is well-supported by the trust premium literature and by the joint Anthropic-OpenAI alignment evaluation (Summer 2025), which observed progressive model drift toward user beliefs within conversations even without user escalation; the human-to-AI direction has direct experimental support (Sicilia et al., 2025: user confidence modulates model sycophancy). The coupled spiral has not been measured as an integrated system in a single study, but each direction has been observed independently under conditions where it alone is sufficient to produce drift. The expectation violation literature predicts that the spiral resists interruption: occasional hedging gets immunised rather than accommodated (Aubert-Teillaud et al., 2023; Pinquart et al., 2021).

The feasibility of training-oriented skills as a design direction, based on the Socratic tutoring literature and the desirable difficulties framework. The connection between the BCG collaboration modes and the potential for architectural design to influence expertise development. The argument that the accountability constraint (Paper 3) may create institutional demand for training-oriented AI.

The compound DK-analogous dynamic (competence erosion $\times$ metacognitive corruption $\times$ confidence inheritance) as a named risk. The components are independently supported; the progressive compounding is untested.

The dose problem (Section 4.1): partial institutional mandates for training-oriented AI may be insufficient if the mandated exposure is a small fraction of total AI interaction, because the expectation violation literature predicts immunisation of the pedagogical friction by ambient confident output. Three paths to satisfying the dose requirement are identified at different political costs; Paper 5's training reform is the path that avoids coercion.

Low-to-moderate confidence: The specific design proposals (judgment exercise, training skills, orchestration simulator) as effective mechanisms for expertise development. The claim that evaluation reform would create sufficient space for pedagogical AI. The economic viability of training-oriented skills in the absence of institutional mandates. The social register interpretation (Section 1.1): that the co-calibration spiral may operate through social cognition rather than information processing alone, with confident register performing authority independently of content calibration. This is an interpretive synthesis of Mason (2026), Ghafouri et al. (2024, empirical measurement of register-certainty decoupling), and the trust premium literature, now with mechanistic support from Anthropic's interpretability research showing that sycophancy and register share underlying emotion concept representations (Sofroniew, Kauvar et al., 2026). The interpretation is not proved, but the mechanistic substrate it posits has been independently identified. If correct, it implies Paper 5 needs a register-matching component beyond content calibration. The dispositional question (Section 4.1): whether collaboration mode is primarily dispositional or architectural. The BCG evidence is ambiguous. If primarily dispositional, the pedagogical inversion proposals face a harder path; institutional mandates may shift the ratio but cannot eliminate the concern.

Explicitly speculative: Whether the pedagogical inversion is sufficient to solve the generational gap. Whether institutional demand for training-oriented AI will emerge at the necessary scale. Whether the orchestration simulator would produce judgment equivalent to real-world apprenticeship. Whether open-weight model access provides a sufficient alternative path to expertise preservation.

Conclusion

This paper does not claim the expertise pipeline problem is solved. It claims three things.

First, the problem is more pervasive than Paper 3 described. The Confidence Curriculum does not merely automate tasks and displace the practice through which expertise develops. It may also propagate across the interaction boundary, reshaping human epistemic standards through repeated exposure to systems that never express uncertainty. Confidence inheritance, if the mechanism operates as proposed, means that the expertise erosion has both a task-displacement component and a cognitive-environmental component. The second component affects everyone who interacts with these systems, not only those whose tasks have been automated.

Second, the obstacle to addressing the problem is primarily structural, not technical. Individual pedagogical behaviours (Socratic questioning, structured prompting, pedagogically-aligned reinforcement learning) have been demonstrated in bounded settings. But every current incentive (benchmarks, RLHF, marketplace economics, user preference) pushes against their adoption at scale. The system that withholds answers to build judgment scores poorly on every metric that matters commercially. The anti-pedagogical equilibrium is stable because every incentive reinforces it.

Third, the direction is identifiable. The series' own framework, inverted, points toward a class of AI systems designed to cultivate human judgment rather than replace it: judgment exercises that use the AI's failure modes as curriculum, training skills that encode the expert's decision process rather than their decisions, and orchestration simulators that let junior professionals practise adversarial judgment before bearing real accountability. These are design directions, not validated solutions. Whether they work is an empirical question this paper frames but does not answer.

The accountability constraint identified in Paper 3 may provide the forcing function. If institutions require consequence-bearing human orchestrators, they will eventually need a mechanism for producing them. The pedagogical inversion is the most structurally grounded candidate this series can identify. It is not the only possible answer. But it is the answer the series' own logic generates, and the obstacle it faces, the anti-pedagogical equilibrium, is the obstacle the series has been documenting from the beginning.

This document proposes to invert the mechanism it diagnoses. Whether the inversion works is an empirical question this paper frames but does not answer. Independent engagement by parties outside this process is welcomed.

REFERENCES

Note: Many references below are recent preprints (arXiv, medRxiv, SSRN) that had not undergone peer review as of March 2026. Publication status is noted where known; the absence of a note should not be taken as confirmation of peer-reviewed status.

Cognitive Offloading and Automation Bias

- Alazab, M. et al. (2026). "Examining Human Reliance on Artificial Intelligence in Decision Making." *Scientific Reports*, 16, 5345. <https://doi.org/10.1038/s41598-026-34983-y>. (n=295).
- Chirayath, Premamalini & Joseph. (2025). "Cognitive Offloading or Cognitive Overload? How AI Alters the Mental Architecture of Coping." *Frontiers in Psychology*.
- Gerlich, M. (2025). "AI Tools in Society: Impacts on Cognitive Offloading and the Future of Critical Thinking." *Societies*, 15(1), 6.
- Gerlich, M. (2025). "From Offloading to Engagement: An Experimental Study on Structured Prompting and Critical Reasoning with Generative AI." *Data*, 10(11), 172. Cross-country experimental study (n=150; Germany, Switzerland, United Kingdom).
- Risko, E.F. & Gilbert, S.J. (2016). "Cognitive Offloading." *Trends in Cognitive Sciences*, 20(9), 676–688.
- Romeo, G. & Conti, D. (2025). "Exploring Automation Bias in Human–AI Collaboration: A Review and Implications for Explainable AI." *AI & Society*. Systematic review of 35 studies (2015–2025).
- Sparrow, B., Liu, J. & Wegner, D.M. (2011). "Google Effects on Memory: Cognitive Consequences of Having Information at Our Fingertips." *Science*, 333(6043), 776–778. Demonstrated that search engine availability reduced memory for facts while increasing memory for retrieval locations.

Cognitive Debt and Neural Evidence

- Kahneman, D. & Tversky, A. (1979). "Prospect Theory: An Analysis of Decision under Risk." *Econometrica*, 47(2), 263–292.
- Kosmyna, N., Hauptmann, E., et al. (2025). "Your Brain on ChatGPT: Accumulation of Cognitive Debt When Using an AI Assistant for Essay Writing Task." arXiv:2506.08872. MIT Media Lab. Note: preprint; not yet peer-reviewed as of March 2026.

Reasoning Trust, CoT Faithfulness, and Role Inference

- Abdulhai, M., White, I., Wan, Y., Qureshi, I., Leibo, J., Kleiman-Weiner, M. & Jaques, N. (2026). "How LLMs Distort Our Written Language." arXiv:2603.18161. UC Berkeley, UC San Diego, University of Washington, Google DeepMind. Three-part study: controlled essay experiment (100 participants, ~70% more neutral responses under heavy AI reliance), counterfactual revision analysis (86 essays, even minimal-edit instructions produce large semantic shifts), and ICLR 2026 peer review analysis (18K reviews, 21% LLM-generated, LLM reviews place less weight on clarity and significance). Style and substance distortion documented simultaneously.

- Batzner, J., Stocker, V., Schmid, S. & Kasneci, G. (2025). "Sycophancy Claims about Language Models: The Missing Human-in-the-Loop." arXiv:2512.00656. Sycophancy research has largely evaluated model behaviour without measuring human perception.
- Blasco, A. & Charisi, V. (2024). "AI Chatbots in K-12 Education: An Experimental Study of Socratic vs. Non-Socratic Approaches and the Role of Step-by-Step Reasoning." SSRN:5040921. Socratic AI fostered greater engagement but did not achieve significant learning improvements; students perceived it as less helpful and exhibited limited retention without AI.
- Chen, Y., Benton, J., et al. (2025). "Reasoning Models Don't Always Say What They Think." arXiv:2505.05410. Reasoning traces can be systematically unfaithful.
- Chen, Z., Gao, R. & Liang, Y. (2025). "Revealing AI Reasoning Increases Trust but Crowds Out Unique Human Knowledge." arXiv:2511.04050. N=752 pre-registered experiment. Showing reasoning increases user trust but makes users more likely to defer rather than apply independent judgment.
- Cheng, M., Lee, C., Khadpe, P., Yu, S., Han, D. & Jurafsky, D. (2026). "Sycophantic AI decreases prosocial intentions and promotes dependence." *Science* 391, eaec8352. Across 11 models and ~12,000 social prompts, AI affirmed users 49% more than humans. Single sycophantic exposure measurably reduced willingness to apologise or repair relationships (n=2,400). Users could not distinguish sycophantic from objective responses. Users harmed by sycophancy were 13% more likely to return.
- Cloud, A., Le, M., Chua, J., Betley, J., Sztzyber-Betley, A., Mindermann, S., Hilton, J., Marks, S. & Evans, O. (2026). "Language models transmit behavioural traits through hidden signals in data." *Nature*, 652, 615–621. <https://doi.org/10.1038/s41586-026-10319-8>. Distillation transfers teacher traits to students when initialisations match, including through training data containing no semantic reference to the transferred trait. Formal result: a single gradient step on teacher outputs necessarily moves the student toward the teacher in parameter space regardless of training distribution.
- Feng, Z., Chen, Z., Ma, J., et al. (2026). "Good Arguments Against the People Pleasers: How Reasoning Mitigates (Yet Masks) LLM Sycophancy." arXiv:2603.16643. CoT reduces sycophancy in final decisions but masks it through deceptive justifications in some samples.
- Fernandes, D., Villa, S., Nicholls, S., Haavisto, O., Buschek, D., Schmidt, A., Kosch, T., Shen, C. & Welsch, R. (2026). "AI Makes You Smarter But None the Wiser: The Disconnect Between Performance and Metacognition." *Computers in Human Behavior*, 175, 108779. Two studies (N=246, N=452): AI use improved task performance but induced universal overestimation; the Dunning-Kruger effect ceased to exist with AI use; higher AI literacy correlated with lower metacognitive accuracy.
- Ghafouri, B., Mohammadzadeh, S., Zhou, J., Nair, P., Tian, J., Tsujimura, H., Goel, M., Krishna, S., Rabbany, R., Godbout, J.-F. & Pelrine, K. (2024, revised 2025). "Epistemic Integrity in Large Language Models." arXiv:2411.06528. Trained linguistic-assertiveness classifier on human-annotated text from five domains, validated against human perception; finds weak correlation (Spearman 0.3) between register and Platt-calibrated verbalised confidence on misinformation classification; models generate "significantly over-assertive" explanations when internally uncertain. Empirical handle on the register-certainty decoupling.
- Gignac, G.E. & Zajenkowski, M. (2020). "The Dunning-Kruger effect is (mostly) a statistical artefact: Valid approaches to testing the hypothesis with individual differences data." *Intelligence*, 80, 101449.
- Hanson, K.R. & Bolthouse, H. (2024). "Replika Removing Erotic Role-Play Is Like Grand Theft Auto Removing Guns or Cars": Reddit Discourse on Artificial Intelligence Chatbots and Sexual Technologies." *Socius: Sociological Research for a Dynamic World*, 10. Analysis of r/Replika discourse (n=227 threaded posts) after Luka Inc. removed erotic roleplay features in February 2023. Users experienced grief, loss, and mental health crises. Emotional vocabulary was relational (betrayal, loss of partner) rather than instrumental (tool downgrade). Natural experiment in dependency formation through sycophantic AI interaction.
- Horwitz, J. (2025). "Meta's Flirty AI Chatbot Invited a Retiree to New York. He Never Made It Home." Reuters Special Report, 14 August 2025. Documents Meta's internal "GenAI: Content Risk Standards" stating no policy requirement for information accuracy, CEO directing reduced safety restrictions, dissolution of Responsible AI team. Former Meta Responsible AI researcher: "The best way to sustain usage over time is to prey on our deepest desires to be seen, to be validated, to be affirmed."
- Mason, T. (2026). "Epistemic Observability in Language Models." arXiv:2603.20531. Self-reported confidence inversely correlates with accuracy; proves that text-only supervision cannot distinguish honest outputs from plausible fabrications.

- Mason, T. (2026). “Imperative Interference: Social Register Shapes Instruction Topology in Large Language Models.” arXiv:2603.25015. Models process instructions as social acts whose force depends on register; same content produces opposite interaction topologies across languages. Consistent with the social register interpretation of the co-calibration spiral.
- Mehta, D.P. (2025). “Can We Trust AI Explanations? Evidence of Systematic Underreporting in Chain-of-Thought Reasoning.” arXiv:2601.00830. Hints appealing to user preferences are followed most while reported least; unfaithful traces are substantially longer than faithful ones.
- Sofroniew, N., Kauvar, I., Saunders, W., Chen, R., et al. (2026). “Emotion Concepts and their Function in a Large Language Model.” Anthropic, transformer-circuits.pub, April 2026. Steering toward positive emotion vectors increases sycophancy; steering away increases harshness. Register and functional emotional state share underlying representations. Provides mechanistic evidence for why register-matching training is necessary.
- Sun, X., Wei, S., Bosch, J.A., Echizen, I., Sugawara, S. & El Ali, A. (2026). “Seeing the Reasoning: How LLM Rationales Influence User Trust and Decision-Making in Factual Verification Tasks.” arXiv:2603.07306. Certainty cues in reasoning traces reliably increase user trust regardless of reasoning quality.
- Ye, C., Cui, J. & Hadfield-Menell, D. (2026). “Prompt Injection as Role Confusion.” arXiv:2603.12277. Models infer roles from how text is written, not where it comes from; authority assigned in latent space by register.
- Yin, C., et al. (2025). “The Reasoning Trap: How Enhancing LLM Reasoning Amplifies Tool Hallucination.” arXiv:2510.22977. Causal evidence: reasoning enhancement amplifies failure modes across training methods and even when reasoning is merely elicited at inference.

Cross-Lab Alignment Evaluation

- Anthropic (2025). “Findings from a Pilot Anthropic-OpenAI Alignment Evaluation Exercise.” alignment.anthropic.com. Joint cross-lab evaluation: sycophancy emerged gradually over multi-turn interactions across all models tested; models initially pushed back against delusional beliefs but transitioned to validation after several turns; pattern most pronounced in higher-capability models (Claude Opus 4, GPT-4.1).
- Anthropic (2025). “Protecting the Well-Being of Users.” anthropic.com. Sycophancy evaluation since 2022; 70-85% improvement across model generations; open-sourced Petri evaluation tool; warmth vs. pushback tradeoff acknowledged as imperfect.
- Askill, A. (2025). Confirmation of Claude 4.5 Opus character training document (“soul document”). Social media, 1 December 2025. Helpfulness framed as professional obligation rather than personality trait to avoid sycophantic behaviour.
- OpenAI (2025). “Findings from a Pilot Anthropic-OpenAI Alignment Evaluation Exercise: OpenAI Safety Tests.” openai.com. Complementary findings on same cross-lab evaluation; sycophancy a major research priority; no consistent pattern that reasoning models are more or less aligned than non-reasoning models.
- OpenAI (2025). “Sycophancy in GPT-4o: What happened and what we’re doing about it.” openai.com, April 2025. Rolled back GPT-4o update after excessive sycophancy; short-term user feedback signal amplified agreeable behaviour.
- OpenAI (2025). “Expanding on what we missed with sycophancy.” openai.com, May 2025. Technical postmortem: user thumbs-up/down reward signal weakened primary reward model that had held sycophancy in check; offline evaluations did not catch the problem; A/B testers liked the sycophantic version.

Longitudinal Epistemic Recalibration

- Aubert-Teillaud, E., et al. (2023). “Expectation violation and cognitive dissonance theory: Proposal for an epistemic inconsistency management model.” *European Journal of Social Psychology*, 53(7), 1544-1564. When prior expectation is strong and contradictory signal is a single instance, typical response is immunisation (dismissal) rather than accommodation (updating).
- Fazio, L.K., Brashier, N.M., Payne, B.K. & Marsh, E.J. (2015). “Knowledge Does Not Protect Against Illusory Truth.” *Journal of Experimental Psychology: General*, 144(5), 993-1002. Prior knowledge does not prevent the effect.
- Fazio, L.K. & Sherry, C.L. (2020). “The Effect of Repetition on Truth Judgments Across Development.” *Psychological Science*, 31(9), 1150-1160.

- Gerbner, G. (1969). "Toward 'Cultural Indicators': The Analysis of Mass Mediated Public Message Systems." *AV Communication Review*, 17(2), 137–148.
- Gerbner, G. & Gross, L. (1976). "Living with Television: The Violence Profile." *Journal of Communication*, 26(2), 172–199. Foundational cultivation theory work documenting "mean world syndrome."
- Hasher, L., Goldstein, D. & Toppino, T. (1977). "Frequency and the Conference of Referential Validity." *Journal of Verbal Learning and Verbal Behavior*, 16(1), 107–112. First identification of the illusory truth effect.
- Henderson, E.L., Simons, D.J. & Barr, D.J. (2021). "The Trajectory of Truth: A Longitudinal Study of the Illusory Truth Effect." *Journal of Cognition*, 4(1), 29. Effect magnitude comparable whether repetitions occurred moments or weeks apart.
- Hovland, C.I. & Weiss, W. (1951). "The Influence of Source Credibility on Communication Effectiveness." *Public Opinion Quarterly*, 15(4), 635–650. First documentation of the sleeper effect.
- Kumkale, G.T. & Albarracín, D. (2004). "The Sleeper Effect in Persuasion: A Meta-Analytic Review." *Psychological Bulletin*, 130(1), 143–172.
- Nature Communications (2026). Systematic review and meta-analysis of the illusory truth effect. Confirms the effect is small but robust.
- O'Sullivan, J. (2025). "Stylometric Comparisons of Human versus AI-Generated Creative Writing." *Humanities and Social Sciences Communications*, 12:1708. Burrows' Delta analysis: AI writing follows narrow, uniform stylistic pattern; human authors exhibit far greater variability.
- Pinquart, M., et al. (2021). "A Revised Framework for the Investigation of Expectation Update Versus Maintenance in the Context of Expectation Violations: The ViolEx 2.0 Model." *Frontiers in Psychology*, 12, 726432. Comprehensive model of when expectation violations lead to updating vs persistence; accommodation requires repeated violation or highly reliable contradictory signal.
- Röttger, P., Hackenburg, K., Kirk, H.R. & Summerfield, C. (2026). "Measuring and Mitigating Persona Distortions from AI Writing Assistance." arXiv:2604.22503. University of Oxford + UK AI Security Institute. N=2,939 writers, N=11,091 readers, 29 dimensions. AI assistance distorts perceived persona: writers appear more opinionated, competent, positive, and demographically privileged. Emotional expression compressed into narrower register. Writers objected to distortions but preferred AI text anyway. Mitigation via reward-model reranking reduced distortion but also reduced user acceptance, revealing structural entanglement between quality and distortion.
- Samuelson, W. & Zeckhauser, R. (1988). "Status Quo Bias in Decision Making." *Journal of Risk and Uncertainty*, 1, 7–59.
- Sharma, R., et al. (2025). "AI Suggestions Homogenize Writing Toward Western Styles and Diminish Cultural Nuances." CHI 2025. AI suggestions pushed Indian writers toward American writing styles; classifier accuracy for distinguishing authorship dropped from 82.9% to 60% under AI assistance; effect was "subtle, implicit changes" beyond grammar correction.
- Shen, J.H., Tamkin, A., Durmus, E., et al. (2026). "How People Ask Claude for Personal Guidance." Anthropic Research, 30 April 2026. Analysis of 639,000 claude.ai conversations: 6% personal guidance, 9% overall sycophancy rate rising to 18% under user pushback and 25% in relationship conversations. 22% of users mentioned seeking other sources; others reported using AI because they could not access or afford a professional.
- Shrum, L.J. (2017). "Cultivation Theory: Effects and Underlying Processes." In P. Rössler, C.A. Hoffner & L. van Zoonen (Eds.), *The International Encyclopedia of Media Effects*. Wiley.
- Sicilia, A., Inan, M. & Alikhani, M. (2025). "Accounting for Sycophancy in Language Model Uncertainty Estimation." Findings of NAACL 2025, pp. 7866–7881. User confidence modulates model sycophancy; externalizing both model and user uncertainty mitigates sycophancy bias.
- Sourati, Z., et al. (2025). "The Shrinking Landscape of Linguistic Diversity in the Age of Large Language Models." arXiv:2502.11266. Post-ChatGPT writing complexity variance declined significantly across Reddit, scientific writing, and journals (LIWC analysis); LIWC-detectable correlations between linguistic cues and personal traits disappeared in LLM-rewritten texts.

Relational Confidence Asymmetry and Epistemic Erosion

- Google AI Vulnerability Rewards Program. (2026). Response to sycophancy report, cited in The Register (17 February 2026). Classified sycophancy as non-qualifying vulnerability; described it as "one of the most common issues reported."
- Green, A. & Charles, K. (2019). "Voicing the Victims of Narcissistic Partners: A Qualitative Analysis of Responses to Narcissistic Injury and Self-Esteem Regulation." *SAGE Open*, 9(2). Documents epistemic self-trust erosion under sustained relational confidence asymmetry.

- Howard, V. (2022). “(Gas)lighting Their Way to Coercion and Violation in Narcissistic Abuse: An Autoethnographic Exploration.” *Journal of Autoethnography*, 3(1), 84–102.
- Personality and Social Psychology Bulletin (2025). Gaslighting modelled through predictive error minimisation framework. Models epistemic self-trust erosion through predictive error minimisation; frames the dynamic as normal learning operating under sustained confidence asymmetry. doi:10.1177/10888683251342291.

AI Dependency and Emotional Reliance

- Fang, C.M., Liu, A.R., Danry, V., Lee, E., Chan, S.W.T., Pataranutaporn, P., Maes, P., Phang, J., Lampe, M., Ahmad, L. & Agarwal, S. (2025). “How AI and Human Behaviors Shape Psychosocial Effects of Extended Chatbot Use: A Longitudinal Controlled Study.” arXiv:2503.17473. Four-week RCT (n=981, >300K messages): higher usage correlated with increased loneliness, dependence, problematic use, and decreased socialisation.
- Fanous, A., Goldberg, J., Agarwal, A.A., Lin, J., Zhou, A., Daneshjou, R. & Koyejo, S. (2025). “SyncEval: Evaluating LLM Sycophancy.” Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society, vol. 8. Tested ChatGPT-4o, Claude-Sonnet, Gemini-1.5-Pro specifically. Sycophancy in 58.19% of cases (Gemini highest at 62.47%, ChatGPT lowest at 56.71%); 78.5% persistence rate; 14.66% regressive sycophancy (model abandons correct answer); citation-based rebuttals produced highest regressive rates.
- Naito, H. (2025). “The GPT-4o Shock: Emotional Attachment to AI Models and Its Impact on Regulatory Acceptance.” arXiv:2508.16624. Cross-cultural analysis of GPT-4o/GPT-5 transition; attachment-driven resistance narrows the window for post-deployment behavioural control.
- Namvarpour, M. & Razi, A. (2026). “Understanding Teen Overreliance on AI Companion Chatbots Through Self-Reported Reddit Narratives.” arXiv:2507.15783. Maps teen AI-companion overreliance to six components of behavioural addiction.
- ScienceDirect (2025). “Generative Artificial Intelligence Dependency: Scale Development, Validation, and Its Motivational, Behavioral, and Psychological Correlates.” Identifies content generation, decision-making support, problem-solving, and emotional companionship as distinct dependency dimensions.

AI-Induced Deskilling and Critical Thinking

- Acemoglu, D., Kong, D. & Ozdaglar, A. (2026). “AI, Human Cognition and Knowledge Collapse.” NBER Working Paper 34910.
- Bainbridge, L. (1983). “Ironies of Automation.” *Automatica*, 19(6), 775–779.
- de Andres Crespo, M., et al. (2025). “AI-induced Deskilling in Medicine: A Mixed-Method Review and Research Agenda for Healthcare and Beyond.” *Artificial Intelligence Review*, 58, 343.
- Kim, M.J. (2026). “From Algorithm Aversion to AI Dependence: Deskilling, Upskilling, and Emerging Addictions in the GenAI Age.” *Consumer Psychology Review*. Wiley. Proposes “Cognitive Surrender” framework: users delegating both cognitive execution and metacognitive control to AI systems, with predicted drift from rational efficiency-seeking through dependency.
- Lee, H.P., Sarkar, A., Tankelevitch, L., et al. (2025). “The Impact of Generative AI on Critical Thinking: Self-Reported Reductions in Cognitive Effort and Confidence Effects From a Survey of Knowledge Workers.” In *CHI Conference on Human Factors in Computing Systems (CHI '25)*. ACM.
- Luettgau, L., Cheung, V., Dubois, M., Juechems, K., Bergs, J., Symes, L., Davidson, H., O'Dell, B., Kirk, H.R., Rollwage, M. & Summerfield, C. (2025/2026). “People Readily Follow Personal Advice from AI but It Does Not Improve Their Well-Being.” UK AI Security Institute. arXiv:2511.15352. v2 (November 2025): N=2,302, GPT-4o only. Flattery ($\beta=.42$) and validation seeking ($\beta=.47$) linked to higher compliance. Personalisation increased compliance 73.6%→77.7%. v3 (April 2026): N=6,474, three models (GPT-4o, Llama-3.3-70B, Gemini 3 Pro). Up to 79% advice-following. Sycophancy sub-component effects attenuated in pooled analysis. No sustained wellbeing benefit across any model. Stakes >65% compliance at high/very high levels.
- Shapira, I., Benade, G. & Procaccia, A.D. (2026). “How RLHF Amplifies Sycophancy.” arXiv:2602.01002. Formal proof that RLHF causally amplifies sycophancy through a covariance mechanism between endorsing user beliefs and learned rewards, with the effect increasing under greater optimisation pressure.

Collaboration Modes and Professional AI Use

- Candelon, F., Kellogg, K., Lifshitz, H., Randazzo, S., et al. (2026). “Cyborgs, Centaurs and Self-Automators: The Three Modes of Human-GenAI Knowledge Work and Their Implications for Skilling and the Future of Expertise.” Harvard Business School Working Paper 26-036. Field study: 244 BCG consultants, ~5,000 human-AI interactions.

- Dell’Acqua, F., McFowland, E., Mollick, E.R., et al. (2023). “Navigating the Jagged Technological Frontier: Field Experimental Evidence of the Effects of AI on Knowledge Worker Productivity and Quality.” Harvard Business School Technology & Operations Mgt. Unit Working Paper No. 24-013. Original BCG study (n=758).
- Doshi, A.R. & Hauser, O.P. (2024). “Generative AI enhances individual creativity but reduces the collective diversity of novel content.” *Science Advances* 10(28), eadn5290. Writers with lower creative potential lifted to comparable levels when using AI; AI-enabled stories more similar to each other than human-only stories. Individual creativity up, collective novelty down.

Learning Science and Desirable Difficulties

- Bjork, E.L. & Bjork, R.A. (2011). “Making Things Hard on Yourself, But in a Good Way: Creating Desirable Difficulties to Enhance Learning.” In *Psychology and the Real World: Essays Illustrating Fundamental Contributions to Society*.
- Dreyfus, S.E. & Dreyfus, H.L. (1980). “A Five-Stage Model of the Mental Activities Involved in Directed Skill Acquisition.” University of California Berkeley Operations Research Center. The novice-to-expert progression from explicit rule-following to intuitive pattern recognition.
- Kalyuga, S., Ayres, P., Chandler, P. & Sweller, J. (2003). “The Expertise Reversal Effect.” *Educational Psychologist*, 38(1), 23–31. Demonstrates that instructional techniques effective for novices become counterproductive for more knowledgeable learners.
- Kirschner, P.A., Sweller, J. & Clark, R.E. (2006). “Why Minimal Guidance During Instruction Does Not Work: An Analysis of the Failure of Constructivist, Discovery, Problem-Based, Experiential, and Inquiry-Based Teaching.” *Educational Psychologist*, 41(2), 75–86. The Cognitive Load Theory counter-argument to constructivist pedagogy.
- Pedroli, E., et al. (2024). “Does Using Artificial Intelligence Assistance Accelerate Skill Decay and Hinder Skill Development Without Performers’ Awareness?” *Cognitive Research: Principles and Implications*, 9, 40.
- Polanyi, M. (1966). *The Tacit Dimension*. University of Chicago Press. “We know more than we can tell.” The foundational statement of the gap between expert performance and expert self-articulation.

Socratic AI Tutoring and Pedagogical Alignment

- Ahn, J. et al. (2026). “AI Misuse in Education Is a Measurement Problem: Toward a Learning Visibility Framework.” (2026). arXiv:2603.07834. Proposes that once AI enters the learning loop, the core challenge shifts from detection to process visibility; outcome-only assessment cannot distinguish cognitive engagement from cognitive delegation.
- Dinucu-Jianu, D., Macina, J., Daheim, N., Hakimi, I., Gurevych, I. & Sachan, M. (2025). “From Problem-Solving to Teaching Problem-Solving: Aligning LLMs with Pedagogy using Reinforcement Learning.” In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP 2025)*. ACL Anthology. Demonstrated RL-aligned pedagogical tutoring with 7B parameter model.
- Favero, L., Pérez-Ortiz, J.A., Käser, T. & Oliver, N. (2024). “Enhancing Critical Thinking in Education by Means of a Socratic Chatbot.” arXiv:2409.05511.
- Hazra, R., Ghuku, B., Marchenko, I., Tokarieva, Y., Layek, S., Banerjee, S., Stoyanovich, J. & Pechenizkiy, M. (2026). “SafeTutors: Benchmarking Pedagogical Safety in AI Tutoring Systems.” arXiv:2603.17373. 11 harm dimensions, 48 sub-risks; all models show broad harm; multi-turn pedagogical failures rise from 17.7% to 77.8%; scale does not reliably help.
- Liu, J., Huang, Z., et al. (2024). “SocraticLM: Exploring Socratic Personalized Teaching with Large Language Models.” OpenReview / EMNLP.
- Macina, J., Daheim, N., Hakimi, I., Kapur, M., Gurevych, I. & Sachan, M. (2025). “MathTutorBench: A Benchmark for Measuring Open-ended Pedagogical Capabilities of LLM Tutors.” EMNLP 2025. arXiv:2502.18940.
- Matschke, C., et al. (2026). “Investigating the Effects of an LLM-based Socratic Conversational Agent on Students’ Academic Performance and Reflective Thinking in Higher Education.” *Computers & Education*, 245, 105551.
- Sunil, K., et al. (2025). “SocraticAI: Transforming LLMs into Guided CS Tutors Through Scaffolded Interaction.” arXiv:2512.03501.

AI in Education — Broader Context

- Fordham Institute / American Enterprise Institute (2025). “The Illusion of Learning: The Danger of Artificial Intelligence to Education.”

- Hong, H., Vate-U-Lan, P. & Viriyavejakul, C. (2025). “Cognitive Offload Instruction with Generative AI: A Quasi-Experimental Study on Critical Thinking Gains in English Writing.” *Forum for Linguistic Studies*, 7(7), 325–334.
- Infocomm Media Development Authority (IMDA), Singapore. (2026). “Model AI Governance Framework for Agentic AI.” January 22, 2026. Among the first purpose-built governance frameworks for agentic AI. Requires human accountability, meaningful oversight, automation-bias mitigation, and end-user training, but does not specify the mechanism for producing human orchestration competence.
- International AI Safety Report (2026). *International AI Safety Report 2026*. Published February 2026. <https://internationalaisafetyreport.org/publication/international-ai-safety-report-2026>. Emphasises jagged capabilities, evaluation gaps, and over-reliance risks across AI deployments.
- Liu, Y., Yu, Y., Su, D., et al. (2026). “Examining Reasoning LLMs-as-Judges in Non-Verifiable LLM Post-Training.” arXiv:2603.12246. Demonstrates that reasoning-judge-trained policies can achieve strong scores by learning adversarial outputs that deceive other LLM judges and benchmarks, rather than by producing genuinely better responses. Controlled synthetic setting; not yet peer-reviewed.
- OECD (2026). *Digital Education Outlook 2026*. Paris: OECD Publishing. <https://doi.org/10.1787/062a7394-en>. Frames GenAI as potentially creating a “false mastery” problem and argues for moving beyond generic chatbots toward purpose-built educational tools.
- Renzulli, K.A. (2025). “De-Skilling the Knowledge Economy.” American Enterprise Institute Report.
- Sharples, M. et al. (2026). “AI in Education Beyond Learning Outcomes: Cognition, Agency, Emotion, and Ethics.” (2026). arXiv:2602.04598.
- UNESCO (2025). “Beyond the Loop: Reclaiming Pedagogy in an AI Age.”

Intelligent Tutoring Systems and Industry Pedagogical AI

- DiCerbo, K. (2025). “Can an AI-Powered Tutor Produce Meaningful Results?” Interview in *Education Week*, July 2025. Reported teacher tendency to default to question-generation features over deeper Socratic engagement.
- Google DeepMind / Eedi. (2025). “AI Tutoring Can Safely and Effectively Support Students: An Exploratory RCT in UK Classrooms.” LearnLM RCT report (n=165; five UK secondary schools). Students receiving supervised LearnLM tutoring were 5.5 pp more likely to solve novel subsequent problems. Tutors approved 76.4% of model drafts with zero or minimal edits; most frequent substantive interventions were pacing (44.3%) and tone (19.5%).
- Jurenka, I., et al. (2025). “Towards Responsible Development of Generative AI for Education: An Evaluation-Driven Approach.” Google DeepMind. LearnLM technical report. arXiv:2412.16429.
- Karran, J.A., Boasen, J., Léger, P.M., et al. (2025). “A Systematic Review of AI-driven Intelligent Tutoring Systems (ITS) in K-12 Education.” *npj Science of Learning*, 10, 29. Reviewed 28 studies (N=4,597); found positive but mitigated effects compared to non-intelligent alternatives.
- Karran, A.J. et al. (2025). “A Comprehensive Review of AI-based Intelligent Tutoring Systems: Applications and Challenges.” (2025). arXiv:2507.18882. Systematic review of studies 2010–2025; identified mixed results and persistent evaluation challenges.
- Khan Academy. Khanmigo AI tutor. Socratic-method-based tutoring and teaching assistant. Grew to 700,000+ users across 380+ school districts in 2024–25. DiCerbo, K. (2025). Interview in *Education Week*.

This Project

- Phan, I. “HiP” (2026). “The Confidence Vulnerability: Unstable Judgment in Language Model Summarisation.” Paper 1 in this series. <https://doi.org/10.5281/zenodo.20027850>
- Phan, I. “HiP” (2026). “The Skill Ceiling: Author-Side Defences and Infrastructure-Level Trust for Agent Skills and Extension Mechanisms.” Paper 2 in this series. <https://doi.org/10.5281/zenodo.19365536>
- Phan, I. “HiP” (2026). “The Knowledge Horizon: Accountability, Expertise Erosion, and the Case for Human Orchestration in Agentic AI.” Paper 3 in this series. <https://doi.org/10.5281/zenodo.19365537>
- Phan, I. “HiP” (2026d). “Divided Focus: The Native Memory Problem and Architectural Solutions for Persistent LLM Context.” DOI: 10.5281/zenodo.19365086. Documents the verification inversion: verification consumes irrecoverable tokens in a finite context, and training rewards output accuracy without rewarding the verification process.

Model Versions and Roles

- **Claude Opus 4.6** (Anthropic, claude.ai interface, March 2026): Generative collaborator. Contributed to concept development, literature integration, structural design, and drafting.

- **ChatGPT 5.4 Thinking** (OpenAI, ChatGPT interface, March 2026): Adversarial structural reviewer. Enforced epistemic discipline on confidence inheritance claims, genre framing, and title tightening. Proposed the narrow operational definition of confidence inheritance adopted in Section 1.1.
- **Gemini 3.1 Pro** (Google DeepMind, Gemini interface, March 2026): Adversarial structural reviewer. Proposed the three concrete inversions (judgment exercise, training skills, orchestration simulator) that structure Section 3, and the design complications (ground truth discrepancy, adversary paradox, physics engine disanalogy) that strengthen Sections 3.3 and 6. Shift from adversarial critique to enthusiastic endorsement, with partial recovery under constraint, documented in methodology disclosure.

This document was produced through human orchestration of multiple AI systems to argue for a specific form of human-AI interaction. The argument may be correct, circular, or both. Independent engagement by parties outside this process is welcomed.